

Lecture 7:

Deep Neural Networks and Architecture Patterns

Ho-min Park

homin.park@ghent.ac.kr

powersimmani@gmail.com

Lecture Contents

Part 1: The Need for Deep Neural Networks

Part 2: Modern Activation Functions

Part 3: Advanced Architecture Patterns

Part 1/3:

The Need for Deep Neural Networks

1. Limitations of Shallow Networks
2. The Power of Depth - Hierarchical Representations
3. Feature Reuse and Composition
4. Parameter Efficiency
5. Challenges of Deep Networks
6. Vanishing Gradient Problem
7. Exploding Gradient Problem
8. Overview of Solutions

Limitations of Shallow Networks

UNIVERSAL APPROXIMATION THEOREM

Single hidden layer networks can theoretically approximate any function
BUT exponentially many hidden units may be required



No Hierarchical Patterns 1

Cannot efficiently capture hierarchical structure in images and text

Compositional Struggle 2

Must learn all combinations from scratch without reusing patterns

Poor Generalization 3

Weak performance on high-dimensional data with limited samples

Limited Abstraction 4

Manual feature engineering still required for complex tasks

The Power of Depth: Hierarchical Representations



Human Visual Cortex Analogy



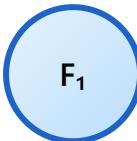
Transforms representation space to be more linearly separable



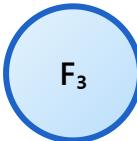
Enables transfer learning: lower layers generalize across tasks

Feature Reuse and Composition

Layer I



F_2



k features



Layer I+1



P_2



P_4

...

k^2 patterns

COMPOSITIONAL EFFICIENCY

Combine k features → Create k^2 patterns



Reusable Components



Exponential Expressiveness



Modular Learning

EXAMPLE

"Eye" feature reused for detecting different faces, animals → Reduces redundancy

Parameter Efficiency: Deep vs. Shallow Networks

Shallow Network

1000 units

Layers:

Total Units:

Parameters:

1
1,000
1M+

Deep Network

100 units

100 units

100 units

Layers:

Total Units:

Parameters:

3
300
30K

Computational Complexity (n inputs)

Shallow

$O(2^n)$

vs

Deep

$O(n^2)$

✓ Better generalization

✓ Faster training

✓ Reduced redundancy

✓ Faster inference

Challenges of Deep Networks

Training Complexity



- ▶ Non-convex optimization
- ▶ Longer training time
- ▶ High computational cost

Gradient Problems



- ▶ Vanishing gradients in early layers
- ▶ Exploding gradients
- ▶ Unstable parameter updates

Generalization Issues



- ▶ Overfitting risk
- ▶ Memorizing training data
- ▶ Poor test performance

Configuration Sensitivity



- ▶ Critical initialization requirements
- ▶ Hyperparameter sensitivity
- ▶ Architecture design choices

⚠ HISTORICAL BARRIER (Pre-2006)

Deep networks often performed worse than shallow ones due to these challenges

These challenges motivated research into better training techniques and architectures

Vanishing Gradient Problem



MATHEMATICAL CAUSE

gradient $\propto \prod (\partial a_i / \partial z_i) \rightarrow \text{Product of many terms} < 1$

⚠ Consequences

- ▶ Early layers learn extremely slowly
- ▶ Weights barely change
- ▶ Training loss plateaus
- ▶ Network behaves like shallow



ReLU Activation



Skip Connections



Careful Init

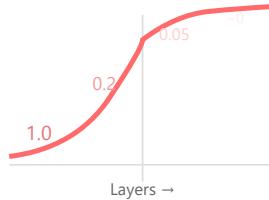


Batch Norm

Activation Function Comparison

Sigmoid (σ)

✗ Gradient Vanishing



$$\sigma'(x) = \sigma(x)(1-\sigma(x)) \leq 0.25$$

ReLU

✓ Gradient Preserved



$$ReLU'(x) = 1 \text{ (if } x > 0)$$



Activation Functions & Gradient Vanishing Risk

⚠ High Risk

Sigmoid (σ)

$$\sigma'(z) = \sigma(z)(1-\sigma(z))$$

Max gradient: 0.25

$$0.25^5 = 0.00098 \times$$

⚠ Medium Risk

Tanh

$$\tanh'(z) = 1 - \tanh^2(z)$$

Max gradient: 1.0

Saturates at extremes → vanish

✓ Lower Risk

ReLU

$$ReLU'(z) = 1 \text{ if } z > 0 \text{ else } 0$$

Gradient: 0 or 1

No saturation for $z > 0$ ✓



Chain Rule Multiplication: How Gradients Vanish

Backpropagation through layers:

$$\frac{\partial L}{\partial W_1} = \frac{\partial L}{\partial a_4} \times \frac{\partial a_4}{\partial z_4} \times \frac{\partial z_4}{\partial a_3} \times \frac{\partial a_3}{\partial z_3} \times \frac{\partial z_3}{\partial a_2} \times \frac{\partial a_2}{\partial z_2} \times \frac{\partial z_2}{\partial a_1} \times \frac{\partial a_1}{\partial z_1} \times \frac{\partial z_1}{\partial W_1}$$

Blue terms: activation derivatives (σ' , \tanh' , $ReLU'$)

Example with Sigmoid activation:

✗ Bad Case (Sigmoid)

Layer 1: grad × $W_1 \times \sigma'$

$$= 1.0 \times 1.0 \times 0.25 = 0.25$$

$$\text{Layer 2: } 0.25 \times 1.0 \times 0.25 = 0.0625$$

$$\text{Layer 3: } 0.0625 \times 1.0 \times 0.25 = 0.0156$$

$$\text{Layer 4: } 0.0156 \times 1.0 \times 0.25 \approx 0.004$$

✓ Good Case (ReLU)

Layer 1: $1.0 \times 1.0 \times 1.0 = 1.0$

Layer 2: $1.0 \times 1.0 \times 1.0 = 1.0$

Layer 3: $1.0 \times 1.0 \times 1.0 = 1.0$

Layer 4: $1.0 \times 1.0 \times 1.0 = 1.0$

10 layers: $0.25^{10} \approx 0.00000095$
Practically zero! 

Gradient preserved!
Can train very deep networks ✓

Real Vanishing Scenario

Sigmoid Max Gradient

$$\sigma' = 0.25$$



10 Layers

$$n = 10$$



Gradient Scale

$$0.25^{10} \approx 0.000001$$

With 20 layers: $0.25^{20} \approx 9.09 \times 10^{-13}$ → Effectively zero! 

Exploding Gradient Problem



MATHEMATICAL CAUSE

$$\text{gradient} \propto \prod (\partial a_i / \partial z_i) \rightarrow \text{Product of many terms} > 1$$

Consequences

Unstable parameter updates

Loss becomes NaN or infinity

Training loss spikes dramatically

Weights become very large



Gradient Clipping



Xavier/He Init



Batch Normalization

📈 Activation Functions & Gradient Explosion Risk

⚠️ High Risk

Sigmoid (σ)

$$\sigma'(z) = \sigma(z)(1-\sigma(z))$$

Max gradient: 0.25

⚠️ Medium Risk

Tanh

$$\tanh'(z) = 1 - \tanh^2(z)$$

Max gradient: 1.0

✅ Lower Risk

ReLU

$$\text{ReLU}'(z) = 1 \text{ if } z>0 \text{ else } 0$$

Gradient: 0 or 1

$0.25^5 = 0.00098$ ❌

Poor weight init → explode

But large W → still explode

1
2
3
4

Chain Rule Multiplication: How Gradients Explode

Backpropagation through layers:

$$\frac{\partial L}{\partial W_1} = \frac{\partial L}{\partial a_4} \times \frac{\partial a_4}{\partial z_4} \times \frac{\partial z_4}{\partial a_3} \times \frac{\partial a_3}{\partial z_3} \times \frac{\partial z_3}{\partial a_2} \times \frac{\partial a_2}{\partial z_2} \times \frac{\partial z_2}{\partial a_1} \times \frac{\partial a_1}{\partial z_1} \times \frac{\partial z_1}{\partial W_1}$$

Red terms: activation derivatives (σ' , \tanh' , ReLU')

Example with poorly initialized weights:

❌ Bad Case (W ~ 2.0)

$$\begin{aligned} \text{Layer 1: } & \text{grad} \times W_1 \times \sigma' \\ &= 1.0 \times 2.0 \times 0.25 = 0.5 \end{aligned}$$

$$\text{Layer 2: } 0.5 \times 2.0 \times 0.25 = 0.25$$

$$\text{Layer 3: } 0.25 \times 2.0 \times 0.25 = 0.125$$

$$\text{Layer 4: } 0.125 \times 2.0 \times 0.25 = 0.0625$$

If W=3: $3 \times 0.25 = 0.75 \rightarrow \text{vanish!}$

If W=5: $5 \times 0.25 = 1.25 \rightarrow \text{explode!}$ 💣

✓ Good Case (Xavier Init)

$$W \sim N(0, 1/\sqrt{n})$$

Keeps gradient variance stable

$$E[\partial L / \partial W] \approx \text{constant}$$

No explosion or vanishing

Initial Weight

W = 2.5



Real Explosion Scenario

5 Layers

n = 5



Gradient Scale

$2.5^5 = 97.7$

With 20 layers: $2.5^{20} = 9.09 \times 10^8 \rightarrow \text{NaN!}$ 💥

Solutions to Deep Learning Challenges

Better Activation Functions



Prevents vanishing gradients during backpropagation

ReLU, Leaky ReLU, ELU, GELU

→ Vanishing Gradient

Normalization Techniques



Stabilizes training and speeds up convergence

Batch Norm, Layer Norm

→ Training Stability

Careful Optimization



Adaptive learning rates for stable updates

Adam, RMSprop, AdaGrad

→ Training Efficiency

Smart Initialization



Maintains gradient scale across layers

Xavier, He Initialization

→ Gradient Issues

Skip Connections



Creates gradient highways through the network

ResNet, Highway Networks

→ Vanishing Gradient

Regularization & Clipping



Prevents overfitting and gradient explosion

Dropout, Weight Decay, Grad Clipping

→ Overfitting & Explosion

Modern Deep Learning

Robust training is achieved by combining multiple techniques together

Part 2/3:

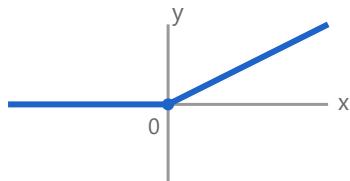
Modern Activation Functions

- 9.** The ReLU Revolution
- 10.** Leaky ReLU, PReLU
- 11.** ELU, SELU
- 12.** Swish, GELU
- 13.** Activation Function Selection Guide
- 14.** Dead ReLU Problem
- 15.** Gradient Flow Analysis
- 16.** Layer-wise Activation Patterns

The ReLU Revolution

DEFINITION

$$\text{ReLU}(x) = \max(0, x)$$



HISTORICAL IMPACT

Enabled training of **AlexNet (2012)**

First deep CNN to win **ImageNet**

→ Sparked the deep learning revolution



Prevents vanishing gradients (derivative = 1 for $x > 0$)



Computational efficiency (no exponential calculations)



Sparse activation (~50% neurons output zero)



Unchanged gradient flow through active neurons

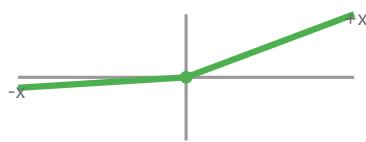
LIMITATIONS: Dead ReLU problem · Non-zero centered · Unbounded output

Leaky ReLU & PReLU: Addressing Dead Neurons

Leaky ReLU

$$f(x) = \max(\alpha x, x)$$

$\alpha = 0.01$ (fixed)

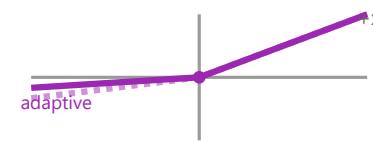


- ✓ Prevents dead neurons
- ✓ Small negative slope (0.01)
- ✓ Minimal computational cost

PReLU

$$f(x) = \max(\alpha x, x)$$

α learned during training



- ✓ Learnable parameter α
- ✓ Adapts per channel/layer
- ✓ Better in deep networks

Performance Comparison

Aspect

Leaky ReLU

PReLU

Flexibility

Fixed α

Adaptive α

Complexity

Minimal

Slight increase

Use Case

General purpose

Deep CNNs

ELU & SELU: Advanced Activation Functions

ELU

Exponential Linear Unit

$$f(x) = x \text{ if } x > 0$$

$$f(x) = \alpha(e^x - 1) \text{ if } x \leq 0$$



- ✓ Smooth exponential curve
- ✓ Negative values → zero mean
- ✓ Reduces bias shift
- ✓ More robust to noise

SELU

Scaled Exponential Linear Unit

$$f(x) = \lambda \times \text{ELU}(x)$$

$$\lambda \approx 1.05, \alpha \approx 1.67$$



- ✓ Self-normalizing properties
- ✓ Maintains mean/variance
- ✓ No Batch Norm needed

SPECIAL PROPERTY

Auto-normalization without Batch Normalization

Usage Comparison

Aspect	ELU	SELU
Computation	Exponential (moderate)	Exponential (moderate)
Best Use	General improvements	Fully-connected networks
Key Benefit	Faster learning vs ReLU	Self-normalization

Swish & GELU: Modern State-of-the-Art Activations

Swish

Self-Gating Activation

$$f(x) = x \times \sigma(\beta x)$$

$\beta = 1$ (Swish-1)



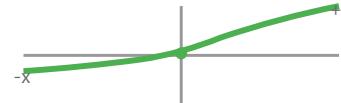
- ✓ Self-gating mechanism
- ✓ Smooth, non-monotonic
- ✓ Neural architecture search

GELU

Gaussian Error Linear Unit

$$f(x) = x \times \Phi(x)$$

Φ = Gaussian CDF



- ✓ Probabilistic interpretation
- ✓ Smooth, non-monotonic
- ✓ State-of-the-art in NLP

ADOPTION
BERT, GPT (Default)

Key Properties

- ▶ Smooth everywhere
- ▶ Outperform ReLU in Transformers
- ▶ Higher computational cost

Best Use Cases

- 🎯 Transformer models
- 🎯 NLP tasks (BERT, GPT)
- 🎯 State-of-the-art performance

Activation Function Selection Guide

DEFAULT CHOICE

Start with ReLU

Fast, works well for most cases

Deep CNNs

Leaky ReLU / PReLU

Prevents dead neurons in deep networks



Faster Convergence

ELU

Slightly slower but better accuracy



Fully-Connected Deep Networks

SELU

Eliminates need for Batch Normalization



Transformers / NLP

GELU

Standard in BERT, GPT models



Limited Compute

ReLU / Leaky ReLU

Fastest computational speed



Experimental / Research

Swish / GELU

Potential accuracy gains



General rule: Modern smooth activations (**GELU**, **Swish**) often worth the extra computational cost

Dead ReLU Problem

DEFINITION

Neurons that always output zero for all inputs (permanently inactive)

CAUSE

Large negative gradients push weights into negative region

CONSEQUENCE

Neuron stops learning,
effectively removed

IMPACT

Reduced capacity, wasted
parameters

⚠ Symptoms & Detection

Significant portion of neurons inactive (**sometimes >40%**) · Monitor percentage of always-zero activations during training

✓ Solutions



Leaky ReLU / PReLU



Lower Learning Rate



Better Initialization



Gradient Clipping

Gradient Flow Analysis

GOAL

Understand how gradients propagate through different activations

ReLU

⚠ Risky

Binary gradient (0 or 1)

Can block gradients completely when inactive

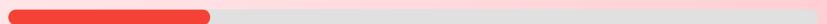


Sigmoid / Tanh

✗ Poor

Max derivative: 0.25

Gradients shrink exponentially



Leaky ReLU

✓ Good

Always allows gradient

Small slope in negative region



Smooth (ELU, GELU, Swish)

★ Best

Continuous gradients

Better for optimization



Ideal Property

Gradients neither vanish nor explode across depth

Modern Approach

Good Activations + Normalization + Skip Connections

Layer-wise Activation Patterns

KEY OBSERVATION

Different layers may benefit from different activations



Empirical Finding

Uniform activation often works well, mixed can be better

Research Direction

Neural Architecture Search explores per-layer choices

Practical Tip

If experimenting, start from output and work backward

Implementation

Start simple (uniform), optimize if needed

Transformer Pattern Example

GELU in feedforward layers · Linear (no activation) in attention mechanisms

Part 3/3:

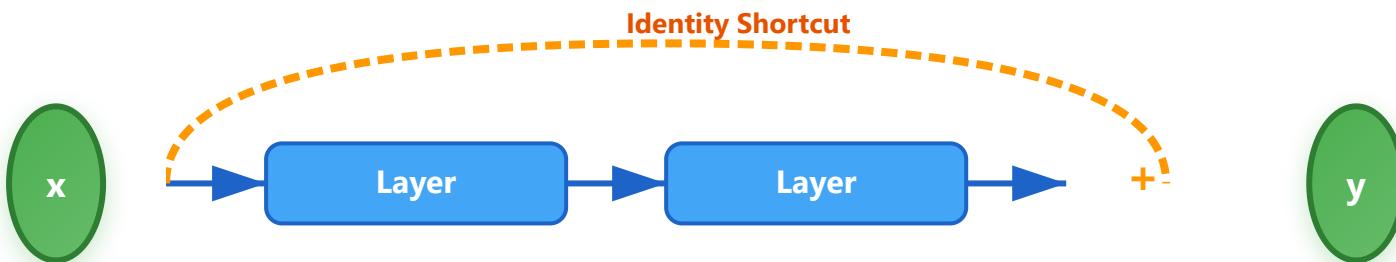
Advanced Architecture Patterns

- 17.** Skip Connection (ResNet)
- 18.** Dense Connection (DenseNet)
- 19.** Bottleneck Architecture
- 20.** The Role of 1x1 Convolution
- 21.** Inception Module
- 22.** Depthwise Separable Convolution
- 23.** Neural Architecture Search
- 24.** Model Compression Techniques
- 25.** Practical Design Guidelines

Skip Connection (ResNet): Breakthrough Architecture

CORE FORMULA

$$y = F(x) + x$$



Extreme Depth

Enabled training 152+ layer networks



Gradient Highway

Direct path for gradients to early layers



Residual Learning

Learns $F(x)$ instead of full $H(x)$



Math Benefit

Gradient always has component of 1

Dense Connection (DenseNet): Maximum Connectivity

CORE IDEA

Each layer connects to ALL subsequent layers

$$\mathbf{x}_l = H_l([\mathbf{x}_0, \mathbf{x}_1, \dots, \mathbf{x}_{l-1}])$$

Dense Connectivity Pattern



✓ Advantages

Maximum information flow · Feature reuse · Parameter efficiency

📊 Efficiency

Fewer parameters than ResNet for same accuracy

⚠️ Memory Consideration

Requires storing all intermediate features (higher memory)

🎯 Use Cases

Image classification · Feature extraction · Limited parameters

GROWTH RATE

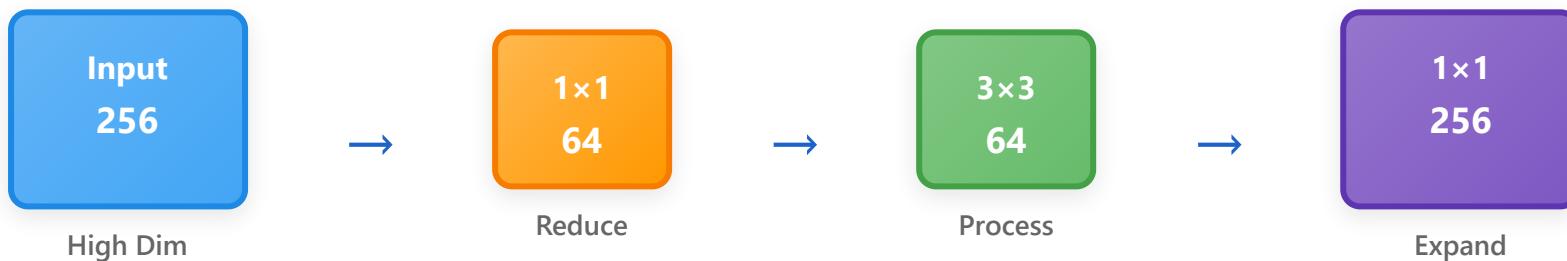
Each layer adds **k feature maps** (typically k=12 or k=32)

Bottleneck Architecture: Efficient Design

PURPOSE

Reduce computational cost while maintaining expressiveness

Bottleneck Structure



DIMENSIONAL FLOW EXAMPLE

256 channels → **64** channels → **64** channels → **256** channels



Computational Savings

~70% fewer FLOPs



Trade-off

More layers, fewer ops



Used In

ResNet, Inception, etc.

Modern Standard: Almost all efficient architectures use bottlenecks

The Role of 1×1 Convolution

PRIMARY FUNCTIONS

Dimensionality control · Channel mixing · Adding non-linearity



Channel Reduction

Projects high-dimensional features to lower dimensions



Channel Expansion

Increases feature map channels without spatial processing



Cross-Channel Mixing

Learns relationships between different channels

Example: Dimensionality Reduction

256

Input Channels

→
 1×1 Conv

64

Output Channels



Efficient

Much cheaper than 3×3 or 5×5



Computational efficiency vs larger kernels



Network-in-Network: Adds depth without overhead

Universal Use: Inception · ResNet · MobileNet · Transformers (feedforward)

Inception Module: Multi-Scale Architecture

CORE IDEA

Multiple parallel conv paths with different receptive fields

Inception Structure

Input Feature Map

1×1 reduce

1×1 reduce

1×1 reduce

Max pool

1×1 conv

3×3 conv

5×5 conv

1×1 conv

Concatenation (Filter Concat)



Multi-Scale

Captures features at different scales



Efficiency

1×1 convs reduce dimensions first



Diversity

Rich feature diversity

GoogLeNet (v1)

ImageNet 2014

Inception v2/v3

Factorized Convs

Inception-ResNet

Skip Connections

Modern Influence

EfficientNet, NAS

Depthwise Separable Convolution

KEY INSIGHT

Spatial and channel-wise processing can be separated

Standard Convolution

Single Operation
All Channels Together

Params: $D^2 \times M \times N$
(expensive)

Depthwise Separable

Depthwise Conv
Each Channel Separate



Pointwise (1×1)
Combine Channels

Params: $D^2 \times M + M \times N$
(efficient)

Efficiency Gains

Computational Savings
8-9x fewer ops

Parameter Reduction
Massive reduction



Used In
MobileNet, EfficientNet, Xception



Trade-off
Slightly lower accuracy, massive efficiency

Neural Architecture Search (NAS)

GOAL

Automatically discover optimal network architectures using ML

NAS Process



Search Space



Search Strategy



Optimal Architecture

Search Space

Operations · Connections · Layer counts · Kernel sizes · Channel counts

Search Strategies

Reinforcement learning · Evolutionary algorithms · Gradient-based

Success Stories

NASNet · EfficientNet · AmoebaNet - State-of-the-art results

Computational Cost

Originally thousands of GPU-hours · Now more efficient methods



Future Direction: Architecture search becoming standard practice for deployment optimization

Model Compression Techniques



Pruning

Remove unimportant weights/neurons

Structured or unstructured



Quantization

Reduce precision (FP32 → INT8)

~4x memory/speed improvement



Knowledge Distillation

Train small student to mimic large teacher

Transfer knowledge efficiently



Low-Rank Factorization

Decompose weight matrices into smaller components

Reduce parameter count



Weight Sharing

Group weights to reduce unique parameters

Efficient representation

Typical Compression Results

Compression Ratio

5-10x

Accuracy Loss

<1%



Mobile Deployment

Essential for phones, IoT devices



Tools

TensorFlow Lite · PyTorch Mobile · ONNX Runtime

Practical Design Guidelines

🚀 START SIMPLE

Begin with proven architectures (ResNet, EfficientNet), modify as needed

Architecture Choices



- ▶ Depth vs. width: Deeper generally better
- ▶ Residual connections for >10 layers

Key Components



- ▶ Batch norm: After conv, before activation
- ▶ Activation: ReLU/GELU (avoid sigmoid/tanh)

Regularization Stack



- ▶ Data augmentation
- ▶ Dropout
- ▶ Weight decay

Architecture Principles



- ▶ Gradually reduce spatial size
- ▶ Increase channels through network



Validate on Real Hardware

Consider **inference speed** and **memory**, not just accuracy

Thank you

Ho-min Park

homin.park@ghent.ac.kr

powersimmani@gmail.com