

Lecture 8:

Loss, Optimization and Scheduling

Ho-min Park

homin.park@ghent.ac.kr

powersimmani@gmail.com

Lecture Contents

Part 1: Loss Function Design

Part 2: Optimization Algorithms

Part 3: Learning Rate Scheduling

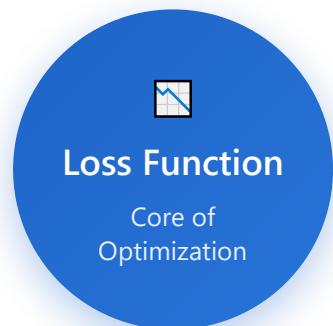
Part 1/3:

Loss Function Design

- 1.** Role and Importance of Loss Functions
- 2.** Regression Losses - MSE, MAE, Huber
- 3.** Classification Loss - Cross-Entropy
- 4.** Hinge Loss and SVM
- 5.** Focal Loss - Class Imbalance
- 6.** Contrastive Loss
- 7.** Triplet Loss
- 8.** Regularization Terms (L1, L2, L1+L2)
- 9.** Custom Loss Function Design

Loss Function: Role & Importance

Understanding the Core Component of Neural Network Training



1 Quantification

Quantifies difference between predicted and actual values



2 Gradient Direction

Guides optimization process by providing gradient direction



3 Task-Specific Design

Different loss functions needed per task (regression vs classification)



4 Performance Impact

Direct impact on model performance and behavior



5 Objective Function

Acts as the objective function that the model seeks to minimize



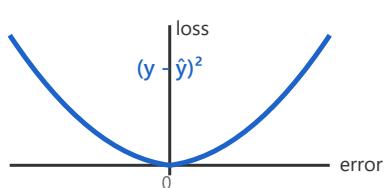
6 Convergence & Accuracy

Affects convergence speed and final model accuracy

Regression Loss Functions Comparison

MSE

Mean Squared Error (L2)



✓ Strong penalty for large errors

✗ Sensitive to outliers

Characteristics from squaring operation

Huber Loss

Hybrid Approach

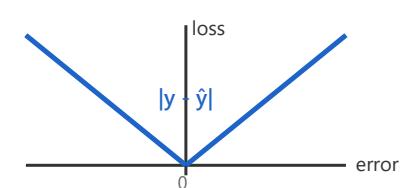


✓ Combines MSE and MAE advantages

✓ Adjustable with delta parameter

MAE

Mean Absolute Error (L1)



✓ Robust to outliers

Uniform penalty regardless of error size

Use MSE when
Outliers are important

Use Huber when
Balanced approach is needed

Use MAE when
Robustness is required



Calculation Examples with 5 Data Points

Sample	Actual (y)	Predicted (\hat{y})	Error (y - \hat{y})
1	100	98	2
2	150	153	-3
3	200	199	1
4	180	182	-2
5	120	140	-20

MSE Calculation

Formula: $MSE = (1/n)\sum(y - \hat{y})^2$

Steps:

- $(2)^2 = 4$
- $(-3)^2 = 9$
- $(1)^2 = 1$
- $(-2)^2 = 4$
- $(-20)^2 = 400$

Sum = 418

$MSE = 418 \div 5$

MSE = 83.6

Huber Loss ($\delta=5$)

Formula:

If $|error| \leq \delta$: $0.5 \times error^2$

If $|error| > \delta$: $\delta \times (|error| - 0.5\delta)$

Steps:

- $|2| \leq 5$: $0.5 \times 4 = 2$
- $|3| \leq 5$: $0.5 \times 9 = 4.5$
- $|1| \leq 5$: $0.5 \times 1 = 0.5$
- $|2| \leq 5$: $0.5 \times 4 = 2$
- $|20| > 5$: $5 \times (20-2.5) = 87.5$

Sum = 96.5

$Huber = 96.5 \div 5$

Huber = 19.3

MAE Calculation

Formula: $MAE = (1/n)\sum|y - \hat{y}|$

Steps:

- $|2| = 2$
- $|-3| = 3$
- $|1| = 1$
- $|-2| = 2$
- $|-20| = 20$

Sum = 28

$MAE = 28 \div 5$

MAE = 5.6

 **Key Insight:**

Notice sample #5 with error of -20 (outlier): **MSE = 83.6** (heavily penalized), **Huber = 19.3** (balanced penalty), **MAE = 5.6** (moderate penalty). Huber Loss provides a middle ground between MSE and MAE!

Classification Loss: Cross-Entropy

Measuring dissimilarity between predicted and true distributions

Core Concept

Measures the difference between predicted probability distribution and actual labels, outputting probability distribution through Softmax activation

- ✓ Encourages high confidence for correct classes
- ✓ Larger penalty for confident wrong predictions
- ≡ Equivalent to maximizing log-likelihood of correct class

1 Binary Cross-Entropy

Used for binary classification problems

2 Classes (0 or 1)

2 Categorical Cross-Entropy

Used for multi-class classification problems

N Classes (Multi-class)



Calculation Examples

Binary Cross-Entropy Examples

Categorical Cross-Entropy Examples

$$\text{Loss} = -[y \cdot \log(p) + (1-y) \cdot \log(1-p)]$$

True=1, Pred=0.95

Loss = 0.051

True=1, Pred=0.70

Loss = 0.357

True=0, Pred=0.10

Loss = 0.105

True=1, Pred=0.50

Loss = 0.693

True=1, Pred=0.20

Loss = 1.609

$$\text{Loss} = -\sum y \cdot \log(p) \text{ (only correct class)}$$

True=[0,1,0], Pred=[0.1,0.8,0.1]

Loss = 0.223

True=[1,0,0], Pred=[0.7,0.2,0.1]

Loss = 0.357

True=[0,0,1], Pred=[0.2,0.2,0.6]

Loss = 0.511

True=[0,1,0], Pred=[0.4,0.4,0.2]

Loss = 0.916

True=[1,0,0], Pred=[0.2,0.5,0.3]

Loss = 1.609

Hinge Loss and Support Vector Machines

Maximum Margin Classification

🎯 Primary Loss Function for SVM

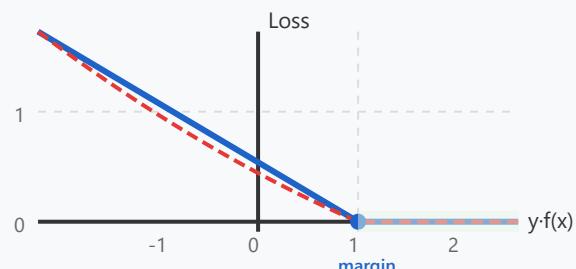
Loss function focused on maximizing the margin between classes

Margin-based Penalty: Penalizes only predictions within margin or misclassified

Confident Predictions: Encourages correct predictions beyond the margin

Outlier Robustness: Less sensitive to outliers compared to cross-entropy

☒ Hinge Loss Graph



— Linear Hinge

— Squared Hinge

L Linear Hinge Loss

$$\max(0, 1 - y \cdot f(x))$$

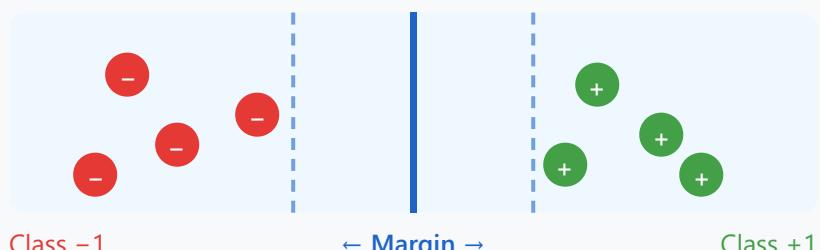
Standard hinge loss with margin-based linear penalty

S Squared Hinge Loss

$$\max(0, 1 - y \cdot f(x))^2$$

Smoother, differentiable variant

🎯 SVM Margin Concept



✓ Use Cases

⚡ Key Advantages

- ✓ Binary classification problems
- ✓ Large-margin classifiers
- ✓ Text classification & document categorization

- ✓ Maximizes decision boundary margin
- ✓ Good generalization on unseen data
- ✓ Sparse solution (support vectors)

Focal Loss: Addressing Class Imbalance

Focusing on Hard Examples

🎯 Target Problem

Designed to solve severe class imbalance problems
Aligns learning with business metrics or task requirements

Focal Loss Formula

$$FL = -\alpha(1-p)^\gamma \log(p)$$

Working Mechanism

- ⬇️ Down-weights loss for well-classified examples
- ⬆️ Focuses learning on hard, misclassified examples
- ✓ Greatly improves minority class performance

📈 Loss Comparison: CE vs Focal Loss

Key Parameters

γ

Focusing Parameter

Controls down-weighting rate
Reduces contribution of well-classified examples

Typical value: $\gamma = 2$

α

Balancing Parameter

Addresses class imbalance problem
Adjusts per-class weights

Addresses class imbalance



How Focal Loss Reweights Samples

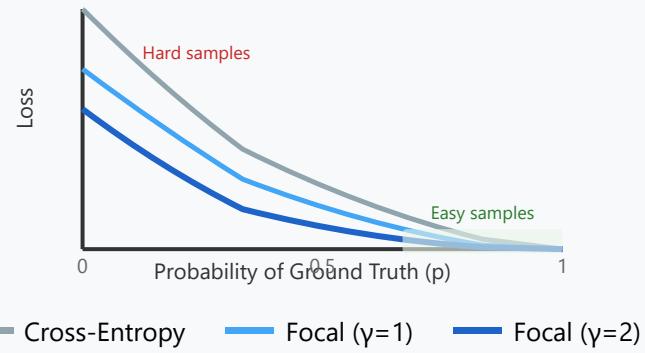
Background
(Easy)



Object
(Hard)



Loss Weight by $(1-p)^\gamma$:



Easy ($p=0.9$) 0.01

Hard ($p=0.2$) 0.64

✓ Easy samples contribute little → Model focuses on hard examples

Contrastive Loss

Learning Similarity Metrics in Embedding Space

Core Concept

Learns meaningful representations in embedding space by training on sample pairs (similar/dissimilar)



Similar Pairs

Place similar samples close together in embedding space

Minimize Distance



Dissimilar Pairs

Place dissimilar samples apart by a margin

Maintain Margin

Training Approach

Can learn meaningful representations without explicit labels

Applications

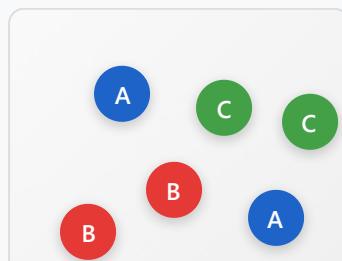
- 1 Face Recognition
- 2 Signature Verification
- 3 Image Retrieval



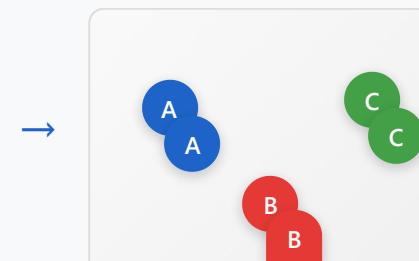
Foundation for Siamese Networks

Embedding Space Transformation

Before Training



After Training



Loss Function Components

$$L = (1-Y) \cdot \frac{1}{2} \cdot D^2 + Y \cdot \frac{1}{2} \cdot \max(0, m - D)^2$$

Y=0 (Similar)
Minimize D^2
Pull together

Y=1 (Dissimilar)
Push if $D < \text{margin}$
Maintain distance

Random

Structured

Similar pairs cluster together • Dissimilar pairs separate



Similar → Close



Dissimilar → Apart

Triplet Loss

Extension of Contrastive Loss with Triplet Structure

Triplet Components

Anchor

Reference sample

Positive

Sample similar to anchor

Negative

Sample different from anchor

Loss Formula

$$\max(d(a,p) - d(a,n) + \text{margin}, 0)$$

Penalizes when anchor is closer to negative than to positive

⚡ Efficiency

More efficient than Contrastive Loss
(One optimization step per triplet)

Triplet Mining

Careful triplet selection strategy needed for effective learning

Hard Negatives: Most difficult negative samples

Semi-Hard Negatives: Medium difficulty negative samples

Mining Approaches

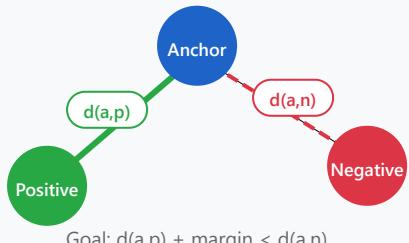
Online Mining

Offline Mining

💡 Applications

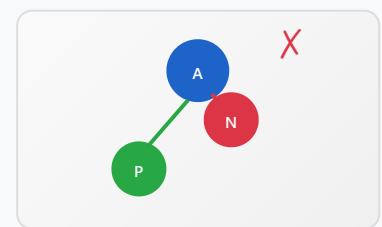
- 1 Face Recognition (FaceNet)
- 2 Person Re-identification
- 3 Fine-grained Similarity Learning

🎯 Triplet Structure in Embedding Space

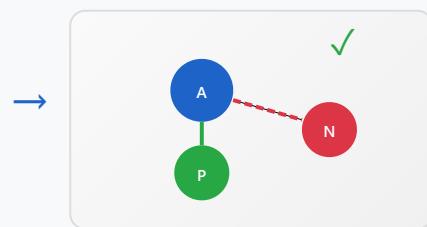


⌚ Learning Process

Before Training



After Training



$d(a,p)$ becomes smaller • $d(a,n)$ becomes larger

✓ Positive should be closer than Negative by at least a margin

Regularization Terms: L1, L2, and Elastic Net

Preventing Overfitting through Model Complexity Constraints



Regularization prevents overfitting by constraining model complexity

L2 Regularization

Ridge

$$\lambda \sum w^2$$

- ✓ Penalizes sum of squared weights
- ✓ Encourages small weights
- ✓ Smooth solutions
- Does not enforce sparsity

L1 Regularization

Lasso

$$\lambda \sum |w|$$

- ✓ Penalizes sum of absolute weights
- ✓ Promotes sparsity
- ✓ Enables automatic feature selection
- Sets some weights to zero

Elastic Net

L1 + L2 Combined

$$\lambda_1 \sum |w| + \lambda_2 \sum w^2$$

- ✓ Combines advantages of L1 and L2
- ✓ Balances sparsity and grouping effect
- ✓ Handles correlated features
- ★ Best of Both Worlds



Lambda (λ): Regularization Strength

Hyperparameter controlling regularization strength - larger values lead to stronger regularization effect

Custom Loss Function Design

Aligning Training with Domain-Specific Objectives

💡 Why Custom Loss?

Standard loss functions may not sufficiently reflect domain-specific objectives

Align learning with business metrics or task requirements

⚙️ Key Considerations

- 1 **Differentiability** - Ensure differentiability
- 2 **Computational Efficiency** - Fast computation
- 3 **Numerical Stability** - Avoid overflow/underflow

⌚ Multi-term Combination

$$L_{\text{total}} = w_1 \cdot L_1 + w_2 \cdot L_2 + \dots + w_n \cdot L_n$$

📋 Design Process

- 1 Verify reflection of true optimization goal
- 2 Combine multiple loss terms with weighting coefficients
- 3 Validate performance improvement with held-out data

🎯 Domain Examples

Computer Vision

Perceptual Loss, SSIM

Object Detection

Classification + Localization + Objectness

Medical Imaging

Sensitivity/Specificity balance

NLP Tasks

Semantic similarity, contextual consistency

🎯 Object Detection Loss Visualization: Success vs Failure Cases

✓ Good Prediction (Low Loss)



✗ Poor Prediction (High Loss)



◻ Ground Truth

◻ Good Prediction

◻ Poor Prediction

① Classification Loss: Cross-entropy for class labels
 $-\sum y \log(\hat{y})$

② Localization Loss: IoU/Smooth L1 for bbox coordinates
Smooth-L1(bbox_pred, bbox_gt)

③ Objectness Loss: Confidence score prediction
BCE(obj_pred, obj_gt)

④ Miss Penalty: False negative detection
Penalizes undetected objects

$$\textbf{Total: } L = \lambda_1 \cdot L_{\text{cls}} + \lambda_2 \cdot L_{\text{box}} + \lambda_3 \cdot L_{\text{obj}}$$

Part 2/3:

Optimization Algorithms

- 10.** Gradient Descent Review
- 11.** Batch vs Mini-Batch vs Stochastic
- 12.** Momentum Method
- 13.** Nesterov Accelerated Gradient
- 14.** AdaGrad
- 15.** RMSprop
- 16.** Adam and AdamW
- 17.** Second-Order Optimization - L-BFGS
- 18.** Comparison of Optimization Algorithms

Gradient Descent Review

Foundation for Modern Neural Network Training

Core Concept

Iterative optimization algorithm to minimize loss function
Updates parameters in negative gradient direction

Update Rule

$$\theta = \theta - \eta \nabla L(\theta)$$

η : Learning Rate

How It Works

- 1 Calculate loss gradient for all parameters
- 2 Move parameters in opposite direction of gradient
- 3 Repeat until convergence

Key Characteristics

Foundation of all modern neural network training algorithms

Convergence depends on learning rate, loss landscape, and initialization

Challenges

- Can get trapped in local minima
- Saddle point problems

Trade-off

Convergence speed \leftrightarrow Stability

Batch vs Mini-Batch vs Stochastic Gradient Descent

Balancing Speed, Stability, and Generalization

Batch GD

Entire Dataset

- ✓ Accurate gradient
- ✓ Stable convergence
- ✗ Slow speed
- ✗ Sharp minima

Mini-Batch GD

Subset of Data

- ✓ Balances speed and stability
- ✓ GPU parallel processing
- ✓ Better generalization
- ✓ Practical choice

Stochastic GD

Single Sample

- ✓ Fast updates
- ✓ Flat minima
- ✗ Noisy gradient
- ✗ Unstable convergence



Typical Mini-Batch Sizes

32

64

128

256

Larger Batches

- Gradients: → More stable
Hardware: → More efficient
Minima: → Sharp

Smaller Batches

- Gradients: → Noisier
Generalization: → Better
Minima: → Flat



Gradient Calculation Examples

Understanding How Each Method Computes Gradients

Example Setup

Dataset: N = 1000 samples | **Loss Function:** $L(\theta) = \frac{1}{2}(y - \hat{y})^2$ | **Learning Rate:** $\alpha = 0.01$

1 Batch Gradient Descent

Gradient:

$$\nabla L(\theta) = (1/N) \sum_{i=1}^N \nabla L(\theta, x_i, y_i)$$

Step 1: Calculate loss for all 1000 samples

Step 2: Compute average gradient: $\nabla L = (1/1000) \sum_{i=1}^{1000} \nabla L_i$

Step 3: Update parameters: $\theta \leftarrow \theta - 0.01 \times \nabla L$

 **Key Feature:** 1 epoch = 1 parameter update (using all 1000 samples)

2 Mini-Batch Gradient Descent

Gradient (batch size B=100):

$$\nabla L(\theta) = (1/B) \sum_{i=1}^B \nabla L(\theta, x_i, y_i)$$

Batch 1: Use samples 1-100 → Compute $\nabla L_1 \rightarrow \theta \leftarrow \theta - 0.01 \times \nabla L_1$

Batch 2: Use samples 101-200 → Compute $\nabla L_2 \rightarrow \theta \leftarrow \theta - 0.01 \times \nabla L_2$

⋮ *Continue...*

Batch 10: Use samples 901-1000 → Compute $\nabla L_{10} \rightarrow \theta \leftarrow \theta - 0.01 \times \nabla L_{10}$

 **Key Feature:** 1 epoch = 10 parameter updates ($1000/100 = 10$ batches)

3 Stochastic Gradient Descent

Gradient (single sample):

$$\nabla L(\theta) = \nabla L(\theta, x_i, y_i)$$

Sample 1: Use $x_1, y_1 \rightarrow$ Compute $\nabla L_1 \rightarrow \theta \leftarrow \theta - 0.01 \times \nabla L_1$

Sample 2: Use $x_2, y_2 \rightarrow$ Compute $\nabla L_2 \rightarrow \theta \leftarrow \theta - 0.01 \times \nabla L_2$

: Continue...

Sample 1000: Use $x_{1000}, y_{1000} \rightarrow$ Compute $\nabla L_{1000} \rightarrow \theta \leftarrow \theta - 0.01 \times \nabla L_{1000}$

 **Key Feature:** 1 epoch = 1000 parameter updates (update after each sample)

Method Comparison Summary

Method	Sample Size	Updates per Epoch	Computation Time	Gradient Accuracy
Batch GD	$N = 1000$	1	Slowest	Very accurate
Mini-Batch GD	$B = 100$	10	Balanced	Sufficiently accurate
Stochastic GD	$B = 1$	1000	Fast updates	Noisy

Momentum Method

Accelerating Gradient Descent with Velocity Accumulation

Core Concept

Accelerates gradient descent by accumulating velocity from past gradients

Update Rules

1. Velocity Update

$$v = \beta v + \nabla L(\theta)$$

2. Parameter Update

$$\theta = \theta - \eta v$$

Momentum Coefficient β : Typically 0.9
(Retains 90% of previous velocity)

Key Benefits

- 1 Helps overcome local minima
- 2 Navigates ravines in loss landscape
- 3 Reduces oscillations in high curvature directions
- 4 Accelerates progress in consistent gradient directions
- 5 Dampens oscillations in narrow valleys



Physical Analogy

Similar to a ball rolling down a hill
gaining momentum

Effect Visualization

Without Momentum

Zigzag movement

With Momentum

Smooth progress

Nesterov Accelerated Gradient (NAG)

Improved Momentum with Look-Ahead Gradient Calculation

Key Innovation

"Look-ahead" mechanism

Calculates gradient at anticipated future position

Update Rules

1. Velocity Update (Look-ahead)

$$v = \beta v + \nabla L(\theta - \beta v)$$

2. Parameter Update

$$\theta = \theta - \eta v$$

Key Point: Calculates gradient at $\theta - \beta v$ position for more accurate direction prediction

vs Standard Momentum

Momentum

Calculates gradient at current position

Nesterov

Calculates gradient at future position

Advantages

- ✓ Better convergence characteristics than standard momentum
- ✓ Reduces overshooting through early trajectory correction
- ✓ Particularly effective for convex optimization problems



Trade-off: Slightly higher computational cost than standard momentum



Widely Used

Widely used in major deep learning frameworks like PyTorch and TensorFlow

AdaGrad (Adaptive Gradient Algorithm)

Parameter-Specific Learning Rate Adaptation

Core Concept

Adaptively adjusts learning rate
for each parameter based on historical gradients

Benefits

- ✓ No manual learning rate tuning needed
- ✓ Advantageous for sparse data

Update Rule

$$\theta = \theta - \eta / \sqrt{G + \epsilon} \odot \nabla L(\theta)$$

G: Cumulative sum of squared gradients

⊗: Element-wise multiplication

Main Drawbacks

- ! Learning rate monotonically decreases over time
- ! Aggressive decay can cause premature training stop



Epsilon Parameter

Prevents division by zero (typically 1e-8)

Adaptation Mechanism

Infrequent Parameters

Larger updates

Frequent Parameters

Smaller updates

Best Use Cases

Sparse Data Processing

Natural Language Processing

RMSprop (Root Mean Square Propagation)

Adaptive Learning Rate with Exponential Moving Average

💡 Key Improvement

Solves AdaGrad's monotonically decreasing learning rate problem

💡 Innovation

Uses exponential moving average of squared gradients

Update Rules

1. EMA of Squared Gradients

$$E[g^2] = \rho E[g^2] + (1-\rho)g^2$$

2. Parameter Update

$$\theta = \theta - \eta / \sqrt{E[g^2] + \epsilon} \odot g$$

vs AdaGrad

AdaGrad

Monotonically decreasing LR

RMSprop

LR can dynamically increase/decrease

⭐ Advantages

- ✓ Learning rate can dynamically increase/decrease
- ✓ Effective for non-stationary objectives
- ✓ Well-suited for Recurrent Neural Networks

📚 Origin

Introduced in Geoff Hinton's Coursera lectures
(Unpublished paper)

⭐ Good default choice for

Decay Rate ρ : Typically 0.9

Maintains moving average of recent gradients

various optimization problems

Adam and AdamW

Most Popular and Improved Optimizers

★ Adam

Most popular optimization algorithm combining Momentum and RMSprop

Moment Estimates

First Moment

Mean

Second Moment

Variance

🎛 Default Parameters

$\beta_1 = 0.9$ (Momentum)

$\beta_2 = 0.999$ (RMSprop)

$\epsilon = 1e-8$ (Numerical stability)

🔧 AdamW

Improved Adam using decoupled weight decay regularization

✨ Key Improvements

- 1 Fixes weight decay implementation bug in original Adam
- 2 Decouples L2 regularization from gradient-based optimization
- 3 Generally better generalization than Adam

Key Difference

Adam

Weight decay in gradient

AdamW

Weight decay decoupled

Includes **Bias Correction**

for moment estimates in early iterations



AdamW is the preferred choice
in modern deep learning

Second-Order Optimization: L-BFGS

Limited-Memory Broyden-Fletcher-Goldfarb-Shanno

Core Concept

Optimization using second-order information (Hessian) with efficient inverse Hessian approximation

Performance Trade-off

Iterations

Fewer ✓

Per-iteration Cost

Higher X

First-Order vs Second-Order

First-Order

- Uses gradient only
- More iterations
- Lower cost per iteration

Second-Order

- Uses Hessian
- Fewer iterations
- Higher cost per iteration

Limitations

- ! Unsuitable for large-scale deep learning due to memory constraints
- ! Not suitable for mini-batch optimization

Suitable Applications

Small Models: Optimization of small-scale models

Scientific Computing: Scientific computing applications

Traditional ML: Logistic Regression, etc.

Characteristics

- ✓ Fewer iterations than first-order methods
- ✓ Deterministic method
- ✗ High computational cost per iteration ($O(n^2)$)

 Requires full batch evaluation

Specific Cases: When full batch evaluation is feasible

Comparison of Optimization Algorithms

Performance Characteristics and Trade-offs

Metric	SGD	Momentum	RMSprop	Adam/AdamW
Training Speed	Slow	Moderate	Fast	Fastest
Generalization	Good	Best	Moderate	Good
Memory Usage	Lowest	Low	Moderate	High
HP Sensitivity	High	Moderate	Low	Lowest

🎬 Optimizer Behavior Visualization



SGD

Slow, zigzag path through loss landscape



[View Animation →](#)



Momentum

Accelerates in consistent direction

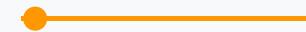


[View Animation →](#)



RMSprop

Adaptive learning per parameter

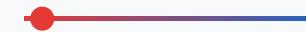


[View Animation →](#)



Adam

Combines momentum + adaptive rates



[View Animation →](#)

SGD + Momentum

Adam/AdamW

RMSprop

Speed Ranking

Adam > RMSprop
> Momentum > SGD

Simple, reliable
Good generalization
Requires tuning

Adaptive
Fast convergence
Out-of-the-box

Good for RNNs
Non-stationary
problems



Recommendation: Start with Adam/AdamW, then fine-tune with SGD+Momentum

[Interactive Visualizations & Resources](#)

[Gradient Descent Viz \(GitHub\)](#)

[Interactive Optimizer Demo](#)

[Why Momentum Works \(Distill\)](#)

[Sebastian Ruder's Guide](#)

Part 3/3:

Learning Rate Scheduling

- 19.** Importance of Learning Rate
- 20.** Fixed vs Adaptive Learning Rate
- 21.** Step Decay
- 22.** Exponential Decay
- 23.** Cosine Annealing
- 24.** Warm-up and Linear Schedule
- 25.** Cyclical Learning Rates

Importance of Learning Rate

The Most Critical Hyperparameter



The most important hyperparameter
affecting training dynamics and convergence

Learning Rate Impact

⚠ Too Large

- Training divergence
- Loss oscillation
- Loss explosion

⚠ Too Small

- Slow convergence
- Stuck in poor local minima

✓ Optimal (Dynamic)

Optimal learning rate changes throughout training

Training Phase Strategy



Early Training

Large learning rate for fast initial progress



Late Training

Small learning rate for fine-tuning and convergence

Key Insights

- 1 Different layers may benefit from different learning rates
- 2 Learning rate scheduling improves final model performance



Visual Comparison: How Learning Rate Affects Optimization

✗ Too Large ($\eta = 1.0$)

✗ Too Small ($\eta = 0.001$)

✓ Optimal ($\eta = 0.1$)



Overshoots the minimum repeatedly, causing oscillation or divergence



Takes tiny steps, extremely slow progress toward the minimum



Balanced steps lead to smooth and efficient convergence

Restart Animation

Fixed vs Adaptive Learning Rate

Comparing Learning Rate Strategies

Fixed Learning Rate

Constant throughout training

- ✓ Simple implementation
- ✗ Suboptimal
- ✗ Requires careful manual tuning
- ✗ Difficult to find good value

Adaptive Learning Rate

Scheduled or metric-based variation

- ✓ Improved convergence
- ✓ Better final accuracy
- ✓ Fast early training
- ✓ Careful late refinement

Adaptive Learning Rate Types

Schedule-based

Predefined decay strategy
(Step, Exponential)

Performance-based

Reduce on validation metric plateau
(Plateau Detection)



Key Benefits of Adaptive LR

- Combines benefits of fast early training and careful late refinement
- Improves both convergence speed and final model performance



Modern Practice: Always use some form of LR Scheduling

Step Decay Learning Rate Scheduling

Reducing Learning Rate at Fixed Intervals

Core Concept

Reduce learning rate by a fixed factor at specific epoch intervals

Formula

$$LR = \text{initial_lr} \times 0.1^{(\text{epoch} // \text{step_size})}$$

Reduce by $0.1 \times$ every `step_size` epochs

Common Schedule

Reduce by $0.1 \times$ every 30 epochs

Epoch 0-29: $LR = 1.0 \times \text{initial_lr}$

Epoch 30-59: $LR = 0.1 \times \text{initial_lr}$

Epoch 60-89: $LR = 0.01 \times \text{initial_lr}$

Characteristics

- ✓ Simple to implement and understand
- ✓ Effective when training duration is known
- ✓ Can align drops with learning regime changes
- ⚠ May cause sudden jumps in loss when LR drops

Learning Rate Progression



Widely Used In

Computer Vision (ResNet Training Schedule)

Exponential Decay Learning Rate Scheduling

Gradual Continuous Reduction

Core Concept

Gradually decrease learning rate following an exponential function

Formula

$$LR = \text{initial_lr} \times \text{decay_rate}^{(\text{epoch} / \text{decay_steps})}$$

Decay Rate: Typically 0.94-0.99
(per epoch or step)

vs Step Decay

Step Decay

Sudden changes
Staircase reduction

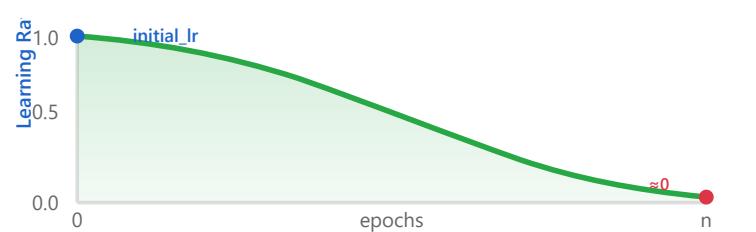
Exponential

Smooth transition
Continuous reduction

Characteristics

- ✓ Smoother than step decay
- ✓ Smooth transition without sudden changes
- ✓ Effective for long training runs
- ⚠ Can be too aggressive if decay rate poorly tuned

Decay Curve Visualization



Commonly Used In

Reinforcement Learning

NLP

Cosine Annealing

Smooth Learning Rate Scheduling with Cosine Function

Core Concept

Learning rate follows a cosine function from maximum to minimum value

Formula

$$LR = \min_{lr} + (\max_{lr} - \min_{lr}) \times (1 + \cos(\pi t/T)) / 2$$

t: current epoch, T: total epochs

Decay Behavior

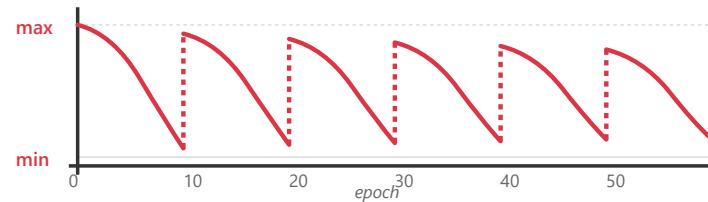
Early Phase

Fast reduction

Late Phase

Slow reduction
(Fine-tuning)

Cosine Curve Visualization



Key Advantages

- ✓ Smooth and gradual decay
- ✓ No hyperparameters to tune (only max/min LR)
- ✓ Improves generalization (empirically proven)

Warm Restarts (SGDR)

Can include restarts to escape local minima

Popular In

Computer Vision

Transformer Training

Warm-up and Linear Schedule

Gradual Start with Progressive Decay

1 Warm-up Phase

- ↑ Gradual increase from low value
- ⌚ Prevents large gradient updates from random initialization
- 📊 Typically first 5-10% of training

2 Linear Decay Phase

- ↓ Linear decrease from peak to minimum
- 🎯 Remaining 90-95% of training

🔗 Combined Schedule

Warm-up → Linear or Cosine Decay

Schedule Visualization



✨ Key Benefits

- ✓ Stabilizes training in early phase
- ✓ Especially effective with large batch sizes
- ✓ Essential for large-scale model training

🎯 Essential For

BERT GPT Vision Transformers LLMs

Cyclical Learning Rates

Oscillating Between Bounds for Better Exploration

Core Concept

Learning rate varies cyclically between minimum and maximum values

Range Policies

Triangular: Linear increase/decrease

Triangular2: Range halves each cycle

Exponential: Exponential variation

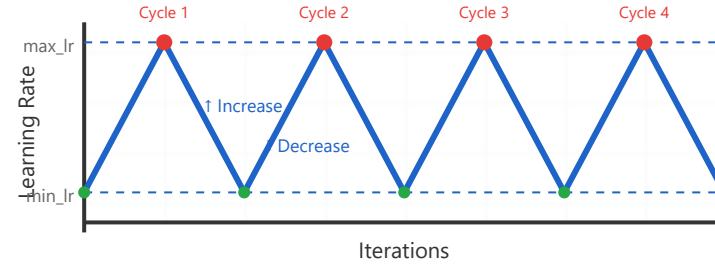
Parameters

Cycle Length: Typically 2-10 epochs per cycle

1cycle Policy

Single cycle including warm-up and cool-down phases

Cyclical Pattern Visualization



Key Benefits

- ✓ Escape saddle points & explore loss landscape
- ✓ Better generalization performance
- ✓ Faster convergence
- ✓ Regularization effect
- ✓ Find better local minima through exploration



Successfully used in
fast.ai training methodology

[Learn more about PyTorch LR Schedulers →](#)

Thank you

Ho-min Park

homin.park@ghent.ac.kr

powersimmani@gmail.com