

Lecture 10:

Data Modality and Feature Extraction

Ho-min Park

homin.park@ghent.ac.kr

powersimmani@gmail.com

Lecture Contents

Part 1: Understanding Data Modalities

Part 2: Traditional Feature Extraction

Part 3: Learning-based Representations

Part 1/3:

Understanding Data Modalities

- 1.** Overview of Data Modalities
- 2.** Structured vs Unstructured Data
- 3.** Text Data Characteristics
- 4.** Image Data Characteristics
- 5.** Audio/Speech Data Characteristics
- 6.** Video Data Characteristics
- 7.** Graph/Network Data
- 8.** Multimodal Data

Overview of Data Modalities

Data modality refers to the type or mode through which information is represented



Text



Images



Audio



Video



Structured Data



Sensor Data



- Unique characteristics per modality

- Specific processing methods needed

- Different storage requirements

- Affects ML algorithm selection

Structured vs Unstructured Data

Structured Data

- Organized in predefined format
- Easy to query and analyze
- Efficient storage
- Standard analytical tools

Examples:

Tables, databases, spreadsheets

VS

Unstructured Data

- No predefined structure
- Rich information content
- Complex processing required
- Advanced analytical methods

Examples:

Text, images, audio, video



Semi-Structured Data

Hybrid format with some organization (JSON, XML)

80-90%

of enterprise data is unstructured and growing rapidly

Structured Data Examples

Classification Example

Classification with PCA Features

Age	Income	Credit Score	Education	Approved
25	45K	680	Bachelor	No
35	75K	720	Master	Yes
42	90K	780	PhD	Yes
28	52K	650	Bachelor	No
50	110K	800	Master	Yes

PC1	PC2	PC3	Gender	Churn
2.4	-1.2	0.8	M	No
-0.5	3.1	-1.5	F	Yes
1.8	0.3	2.2	M	No
-2.1	2.7	-0.9	F	Yes
3.2	-0.8	1.4	M	No

Regression Example

Size (sq ft)	Bedrooms	Age (years)	Distance (km)	Price (\$K)
1200	2	10	5.2	320
1800	3	5	3.1	485
2400	4	2	1.8	650
950	1	15	8.5	245
2100	3	8	4.3	560

Regression with PCA Features

PC1	PC2	PC3	PC4	Sales (\$M)
12.5	-3.2	1.8	0.5	2.4
8.7	4.1	-2.3	1.2	3.8
15.3	-1.5	3.4	-0.8	5.2
6.2	2.8	-1.1	0.9	1.9
11.8	0.3	2.7	-0.4	4.5

Text Data Characteristics



Sequential Nature

Word order and context matter significantly



High Dimensionality

Vocabulary size can reach tens of thousands



Sparse Representation

Most words don't appear in any given document



Ambiguity

Words have multiple meanings depending on context



Language-Dependent

Different languages have unique structures and rules



Rich Semantic Info

Information encoded in syntax and grammar



Required Processing Steps

Tokenization

Normalization

Vocabulary Management



Text Data Examples

Example 1: Product Review

"This laptop is AMAZING!!! Best purchase ever 😊"

Example 2: Social Media Post

"Can't believe it's already 2024... time flies! #newyear"

Example 3: Customer Inquiry

"Hi, I'm wondering if this product comes in blue?"

🔧 Processing Steps Example

1 Tokenization

Split text into individual tokens (words, subwords, or characters)

Input: "This laptop is AMAZING!!!"

Output: ["This", "laptop", "is", "AMAZING", "!", "!", "!"]

2 Normalization

Convert to lowercase, remove punctuation, handle special characters

Input: ["This", "laptop", "is", "AMAZING", "!", "!", "!"]

Output: ["this", "laptop", "is", "amazing"]

3 Vocabulary Management

Map tokens to unique IDs based on vocabulary dictionary

Input: ["this", "laptop", "is", "amazing"]

Vocabulary: {this: 45, laptop: 1203, is: 12, amazing: 789}

Output: [45, 1203, 12, 789]

Image Data Characteristics



Spatial Structure

Pixel relationships and local patterns are critical



High Dimensionality

Small images contain thousands of features



Translation Invariance

Objects recognized regardless of position



Scale Variance

Same object can appear at different sizes



Color Channels

RGB, grayscale, or other color spaces



Resource Intensive

Significant storage and computation needed

Hierarchical Features

Edges



Textures



Parts



Objects



Requires significant storage and computational resources

Audio/Speech Data Characteristics



Temporal Structure

Time-dependent sequential information



Frequency Domain

Sound is combination of different frequencies



Sampling Rate

Typically 16kHz-44.1kHz for speech and music



Variable Length

Utterances and clips have different durations



Background Noise

Acoustic environment affects quality



Speaker Variability

Accent, pitch, speed vary across individuals



Required Time-Frequency Analysis

Spectrograms

MFCCs

Mel-Frequency Analysis

Video Data Characteristics

Spatial (Image)

Temporal (Sequence)

Video



High Dimensionality

Multiple frames per second



Motion Patterns

Object & camera motion



Temporal Redundancy

Similar consecutive frames



Frame Rate

24-60 fps typical for video



Resource Requirements

Massive storage & processing power



Optimization Strategy

Can be decomposed into keyframes for efficient processing

⌚ How Video is Composed

Frame 1



Frame 2



Frame 3



Frame 4



...

 **Key Concept:** Video = Sequence of frames where keyframes store complete information, and intermediate frames only store changes (deltas) from previous frames

Graph/Network Data

Represents relationships and connections between entities



Nodes (Vertices)

Represent objects or entities



Edges

Relationships between nodes

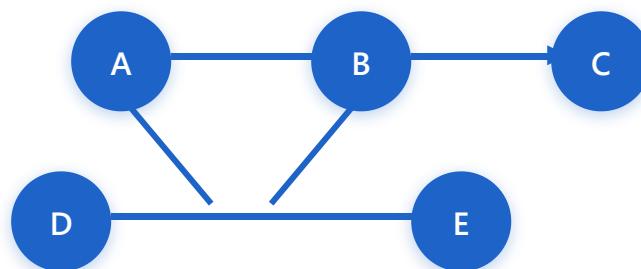
Directed

Undirected

Weighted

Unweighted

Example Graph Structure



Common Applications

Social Networks

Molecular Structures

Knowledge Graphs

Non-Euclidean Structure → Requires Specialized Algorithms:

Graph Neural Networks (GNNs)

Graph Feature Extraction Process



Step-by-Step Feature Extraction

1

Graph Representation

Convert graph to mathematical structures (adjacency matrix, edge list)

2

Feature Computation

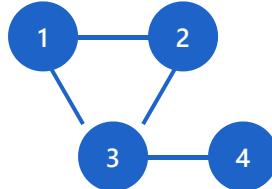
Calculate node-level and graph-level features

3

Feature Vector

Aggregate features into numerical vectors for ML models

Sample Graph



Adjacency Matrix

	1	2	3	4
1	0	1	1	0
2	1	0	0	1
3	1	0	0	1
4	0	1	1	0

Types of Graph Features

◆ Node-Level Features

Degree: Number of connections (노드 1: degree = 2)

Centrality: Importance in the network

Clustering Coefficient: How connected neighbors are

PageRank: Node influence score

◆ Edge-Level Features

Weight: Connection strength

Distance: Shortest path length

Common Neighbors: Shared connections

◆ Graph-Level Features

Number of Nodes: Total vertices (예시: 4개)

Number of Edges: Total connections (예시: 4개)

Density: Connectivity ratio

Diameter: Maximum distance between nodes

◆ Structural Features

Triangles: Number of 3-node cycles

Connected Components: Separate subgraphs

Average Path Length: Mean distance between nodes



Key Insight: These features transform non-Euclidean graph data into numerical vectors that machine learning models can process!

Practical Example: Social Network Analysis



Friend Network Feature Extraction

```
# Python example using NetworkX
import networkx as nx
import numpy as np

# 1. Create graph
G = nx.Graph()
G.add_edges_from([(1,2), (1,3), (2,4), (3,4)]) # 2. Extract node features
degrees = dict(G.degree())
{1:2, 2:2, 3:2, 4:2}
centrality = nx.betweenness_centrality(G)
clustering = nx.clustering(G) # 3. Extract graph features
num_nodes = G.number_of_nodes() # 4 num_edges = G.number_of_edges() # 4 density =
nx.density(G) # 0.667 avg_degree = np.mean(list(degrees.values())) # 2.0 # 4. Create feature vector for Node 1
node_1_features = [degrees[1], # degree: 2
centrality[1], # centrality: 0.0
clustering[1] # clustering: 1.0 ] # Result: [2, 0.0, 1.0]
```



Graph Neural Network (GNN) Process

1

Initial Features

Each node starts with feature vector
(e.g., [2, 0.0, 1.0])

2

Message Passing

Nodes aggregate features from
neighbors

3

Feature Update

Neural network combines local +
neighbor features

```
# GNN-style feature aggregation (simplified)
def aggregate_features(node, neighbors, features):
    # Gather neighbor features
    neighbor_features = [features[n] for n in neighbors]
    # Aggregate (e.g., mean pooling)
    aggregated = np.mean(neighbor_features, axis=0)
    # Combine with node's own features
    updated = neural_network([features[node], aggregated])
    return updated
# For Node 1 with neighbors [2, 3]:
# Input: own features + aggregated neighbor features
# Output: enriched representation capturing local graph structure
```

Feature Extraction Pipeline Summary



Traditional ML Approach

1. Manual Feature Engineering

- Calculate statistical features
- Create handcrafted features
- Limited to known patterns

2. Fixed Representation

- Features don't adapt
- Same for all tasks



GNN Approach

1. Learned Representations

- Neural networks learn features
- Captures complex patterns
- Discovers hidden structures

2. Task-Specific

- Adapts to specific problems
- End-to-end learning



Key Takeaway

Graph feature extraction transforms complex relational data into numerical vectors
that preserve structural information for machine learning tasks.

GNNs automate and optimize this process through learned representations!

Multimodal Data

Combines information from multiple modalities simultaneously



Video with Audio



Image with Captions



Sensor with Location



Benefits

- Complementary information beyond single modality
- Improved robustness: other modalities compensate for noise
- Richer understanding of data



Challenges

- Alignment across modalities
- Synchronization of data streams
- Fusion of different modalities

Requirements

Cross-Modal Learning

Joint Representations

Applications

Video Understanding

Visual Question Answering

Robotics

Part 2/3:

Traditional Feature Extraction

- 9.** Feature Engineering Principles
- 10.** Text - BoW, TF-IDF
- 11.** Text - N-gram, POS
- 12.** Image - Edge, Corner Detection
- 13.** Image - SIFT, SURF, HOG
- 14.** Audio - FFT, Spectrogram
- 15.** Audio - MFCC, Chroma
- 16.** Time Series - Statistical Features

Feature Engineering Principles

Features are measurable properties used to represent data for ML models

Good Features Are

✓ Informative

✓ Discriminative

✓ Independent



Domain knowledge is crucial for effective feature engineering



Feature quality often more important than model complexity



Consider relevance, redundancy, and computational cost



Feature scaling and normalization improve model performance



Feature quality often more important than model complexity

Iterative Process

Extract



Evaluate



Refine



Text Feature Extraction: BoW & TF-IDF



Bag of Words (BoW)

Represents text as word frequency counts

- Ignores word order and grammar
- Only considers word occurrence
- Simple frequency-based representation



TF-IDF

Term Frequency-Inverse Document Frequency weighting

- Balances local and global importance
- Reduces weight of common words
- Increases weight of rare words

TF-IDF Components

TF (Term Frequency)

How often word appears in document (local importance)



IDF (Inverse Document Frequency)

How rare word is across documents (global importance)

✓ **Advantages:** Simple and interpretable

✗ **Limitations:** Loses semantic and syntactic information

Bag of Words (BoW) - Practical Example

Example Documents:

Doc 1: "I love machine learning"

Doc 2: "Machine learning is amazing"

Doc 3: "I love deep learning"

Step 1: Build Vocabulary

Unique words across all documents:

I love machine learning is amazing deep

Step 2: Count Word Frequencies

Document	I	love	machine	learning	is	amazing	deep
Doc 1	1	1	1	1	0	0	0
Doc 2	0	0	1	1	1	1	0
Doc 3	1	1	0	1	0	0	1

 **Result Interpretation:** Each document is represented as a 7-dimensional vector. Example: Doc 1 = [1, 1, 1, 1, 0, 0, 0]
→ Word order and grammar are ignored; only word occurrence counts are considered.

TF-IDF - Practical Example

Same Example Documents:

Doc 1: "I love machine learning"

Doc 2: "Machine learning is amazing"

Doc 3: "I love deep learning"

TF-IDF Formula:

$$\text{TF-IDF}(\text{word}, \text{doc}) = \text{TF}(\text{word}, \text{doc}) \times \text{IDF}(\text{word})$$

$$\text{TF}(\text{word}, \text{doc}) = (\text{word count in doc}) / (\text{total words in doc})$$

$$\text{IDF}(\text{word}) = \log(\text{total documents} / \text{documents containing word})$$

Example: Calculate TF-IDF for "learning" in Doc 1

Step 1: Calculate TF

$$TF("learning", Doc 1) = 1 / 4 = 0.25$$

(appears 1 time / total 4 words)

Step 2: Calculate IDF

$$IDF("learning") = \log(3 / 3) = \log(1) = 0$$

(3 documents / 3 documents containing "learning")

Step 3: Calculate TF-IDF

$$TF-IDF("learning", Doc 1) = 0.25 \times 0 = 0$$

→ "learning" appears in all documents, so it has low importance!

Example: Calculate TF-IDF for "deep" in Doc 3

Step 1: Calculate TF

$$TF("deep", Doc 3) = 1 / 4 = 0.25$$

Step 2: Calculate IDF

$$IDF("deep") = \log(3 / 1) = \log(3) \approx 1.099$$

("deep" appears in only 1 document → rare word)

Step 3: Calculate TF-IDF

$$TF-IDF("deep", Doc 3) = 0.25 \times 1.099 \approx 0.275$$

→ "deep" is a rare word, so it has high importance!

 **Key Point:** TF-IDF reduces the weight of common words (e.g., "learning") and increases the weight of rare words (e.g., "deep") to better represent document characteristics.

Text Feature Extraction: N-grams & POS



N-grams

Contiguous sequences of N words

Examples:

Bigrams (2): "machine learning"
Trigrams (3): "natural language processing"

- ✓ Capture local word order
- ✓ Identify common phrases
- ✓ Provide more context than single words



POS Tagging

Identifies grammatical roles of words

Tags:

Noun, Verb, Adjective, Adverb...

- ✓ Identify syntactic patterns
- ✓ Analyze sentence structure
- ✓ Distinguish word usage contexts



Combined Benefits

Using N-grams and POS tags together improves text classification and information extraction



Trade-off

Between context capture and computational complexity (N-grams increase feature space)



Practical Processing Examples

1 Input Sentence:

"I love machine learning"



2 N-gram Extraction:

Unigrams (1):

["I", "love", "machine", "learning"]

Bigrams (2):

["I love", "love machine", "machine learning"]

Trigrams (3):

["I love machine", "love machine learning"]

Features Generated:

Feature Vector:

[1, 1, 1, 1, 1, 1, 1, ...]

(Total: 4 unigrams + 3 bigrams + 2 trigrams)

1 Input Sentence:

"The quick brown fox jumps"



2 POS Tagging:

The/DET quick/ADJ brown/ADJ
fox/NOUN jumps/VERB

Tags Extracted:

- DET: Determiner (1)
- ADJ: Adjective (2)
- NOUN: Noun (1)
- VERB: Verb (1)



Features Generated:

POS Pattern:

"DET-ADJ-ADJ-NOUN-VERB"

POS Counts:

[DET:1, ADJ:2, NOUN:1, VERB:1]

Image Feature Extraction: Edge & Corner Detection



Edge Detection

Boundaries where intensity changes sharply

Detection Operators:

Sobel

Prewitt

Canny



Corner Detection

Points where edges intersect or change direction rapidly

Detection Method:

Harris Corner Detector: Identifies distinctive points for matching



Feature Hierarchy

Low-level features (edges, corners) are fundamental for higher-level understanding

Foundation for object detection and image matching



Robust to illumination changes



Resilient to minor transformations



Visual Examples



Edge Detection - Building Example

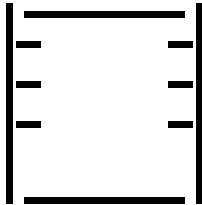
Original Image

Sobel/Canny Edge

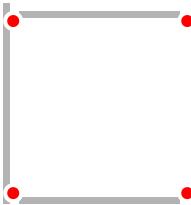
Harris Corner



Full color building with texture and details



Outlines extracted - object boundaries detected



Key points identified - useful for matching



Use Cases

Object Recognition



Edge detection helps identify object boundaries for classification and recognition

Image Matching



Corner points serve as landmarks for matching images across different views

Motion Tracking



Corners are stable features that can be tracked across video frames

Image Descriptors: SIFT, SURF & HOG



SIFT

Scale-Invariant Feature Transform

Detects and describes local features

- ✓ Scale invariant
- ✓ Rotation invariant
- ✓ Partial illumination invariance



SURF

Speeded Up Robust Features

Faster approximation of SIFT

- ✓ Improved speed
- ✓ Similar robustness
- ✓ Efficient computation



HOG

Histogram of Oriented Gradients

Captures edge orientation distribution

- ✓ Shape description
- ✓ Edge patterns
- ✓ Object detection

Capabilities

Enable robust object recognition and image matching



Historical Context: Widely used before deep learning era for computer vision tasks



Current Use: Still valuable for lightweight applications and limited data scenarios

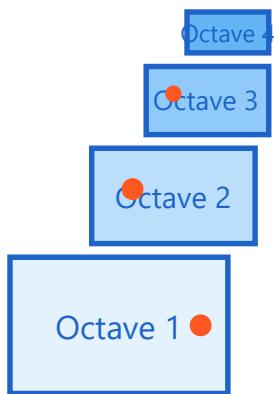


SIFT: Scale-Invariant Feature Transform

💡 Core Principle

SIFT finds feature points in **scale space** and generates descriptors using **gradient orientation histograms** around each keypoint.

DoG Pyramid (Difference of Gaussians)



Multiple Scales

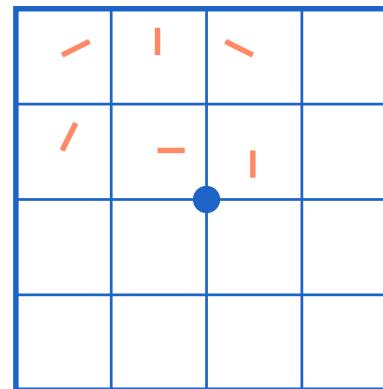
Apply Gaussian blur at each scale

● Keypoints

Detected at extrema locations

Divide image into multiple scales and compute DoG to find scale-invariant features

SIFT Descriptor (128-dimensional)



$4 \times 4 = 16$ blocks

Each block:
8-bin histogram

$16 \times 8 = 128$
dimensional vector

Divide 16x16 region into 4x4 blocks and compute 8-direction gradient histogram per block

1

Scale-Space Extrema Detection

Find extrema in DoG pyramid → Detect scale-invariant keypoint candidates



2

Keypoint Localization

Refine with Taylor expansion → Remove unstable points



3

Orientation Assignment

Gradient orientation histogram → Assign dominant orientation (rotation invariance)

**4**

Descriptor Generation

$4 \times 4 \text{ blocks} \times 8 \text{ orientations} = 128\text{-dimensional feature vector}$

SURF: Speeded Up Robust Features

Core Principle

SURF uses **Integral Image** and **Box Filters** to approximate SIFT faster. Detects keypoints based on Hessian matrix.

Integral Image

Original Image
3 5 2 ...
1 7 4 ...
6 2 8 ...

Trans

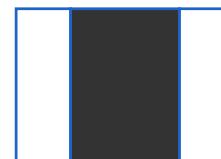
Integral Image
3 8 10 ...
4 16 22 ...
10 24 38 ...

Fast Region Sum

Calculate sum of any rectangle in 4 operations

With Integral Image, box filters can be computed in $O(1)$ time

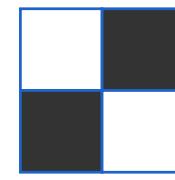
Fast-Hessian Detector



D_{xx}



D_{yy}



D_{xy}

Approximate Hessian with Box Filters

$$\text{Det}(H) \approx D_{xx} \cdot D_{yy} - (0.9 \cdot D_{xy})^2$$

Max value location = Keypoint

Quickly approximate Hessian matrix with 3 box filters to detect blobs

1

Generate Integral Image

Original image → Integral Image (preprocessing)



2

Fast-Hessian Detector

Calculate Hessian determinant with box filters → Detect keypoints



Orientation via Haar Wavelet

3 Determine dominant orientation using Haar wavelet responses



4 **64D Descriptor**

4×4 sub-regions \times 4D Haar response = 64-dimensional vector

⚡ **Speed Improvement:** 2-3x faster than SIFT, with Integral Image enabling same speed for all scales



HOG: Histogram of Oriented Gradients

💡 Core Principle

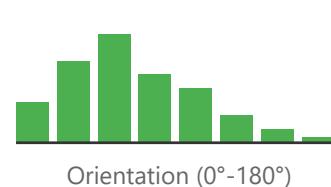
HOG divides images into small **cells** and computes **gradient orientation histograms** in each cell. Normalizes by blocks to create illumination-robust shape descriptors.

Gradient Computation & Orientation



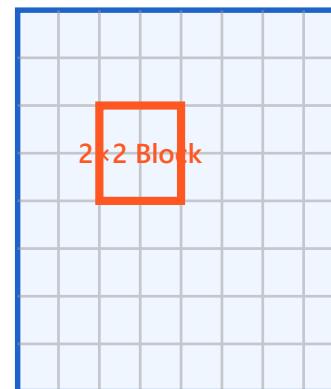
Aggregate gradient orientations in each cell (8×8) to 9 bins

9-bin Histogram



Calculate gradient magnitude and direction at each pixel and aggregate into histogram

Cell & Block Structure



64×128 Window

Group 2 \times 2 cells into blocks and normalize to make robust against illumination changes

Structure:

- Cell: 8×8 pixels
- Block: 2 \times 2 cells
- Stride: 8 pixels

Dimension:

7 blocks (H) \times 15 blocks (V)
 \times 4 cells \times 9 orientations
= 3,780 dimensions

1

Gradient Computation

Calculate x, y direction gradients at each pixel (using [-1,0,1] filter)



2

Orientation Binning

Divide into 8×8 cells, generate 9-bin orientation histogram in each cell



3**Block Normalization**

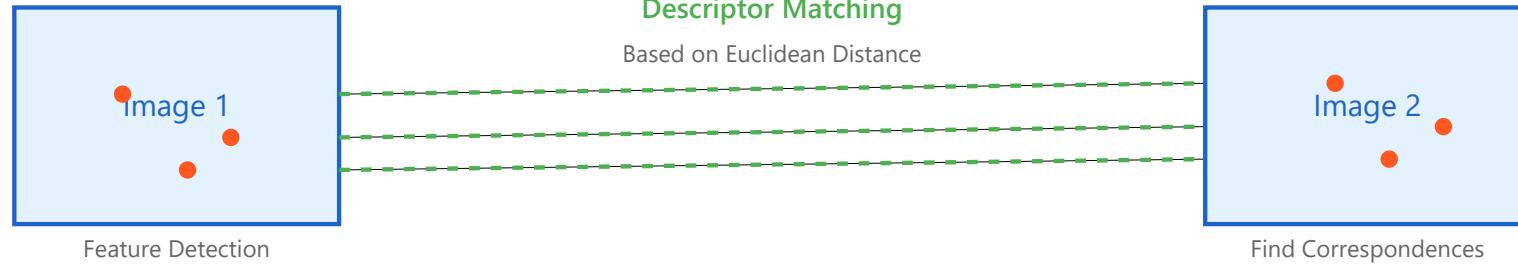
Group into 2×2 cell blocks and apply L2 normalization (50% overlap)

**4****Feature Vector**

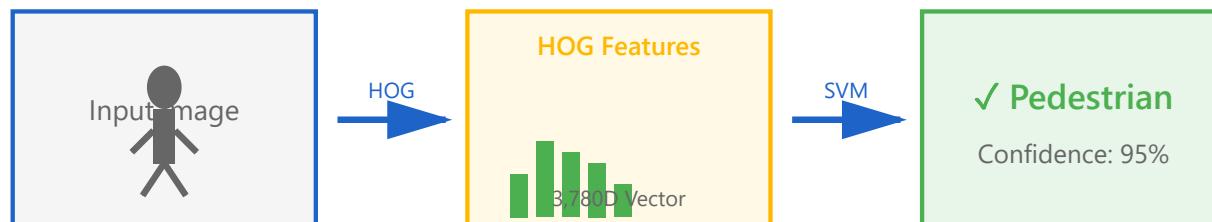
Concatenate all block histograms \rightarrow 3,780-dimensional descriptor

Real-world Applications

Image Matching (SIFT/SURF)



Object Detection (HOG + SVM)



Extract HOG features from sliding windows and detect pedestrians with pre-trained SVM

Detailed Algorithm Comparison

Characteristic	SIFT	SURF	HOG
Core Method	DoG (Difference of Gaussians)	Fast-Hessian (Box Filter)	Gradient Orientation Histogram
Feature Detection	Scale-space extrema	Hessian determinant maxima	Dense sampling (entire window)
Descriptor Dimension	128-dimensional	64-D (standard) / 128-D	3,780-D (64x128 window)
Computation Speed	Slow (baseline)	Fast (2-3x faster than SIFT)	Medium (depends on window size)
Scale Invariance	Excellent	Excellent	Limited (fixed window)
Rotation Invariance	Excellent	Excellent	Limited
Main Applications	Image matching, panorama, 3D reconstruction	Real-time tracking, video processing	Pedestrian detection, object recognition
Patent Status	Expired (2020)	Patent protected	Free to use

Selection Guide

Need high accuracy → SIFT (image stitching, 3D reconstruction, object recognition)

Need real-time processing → SURF (video tracking, AR applications, real-time matching)

Object detection goal → HOG + SVM (pedestrian, vehicle, face detection)



Key Difference: SIFT/SURF are keypoint-based, HOG is dense descriptor. Choose based on your application.



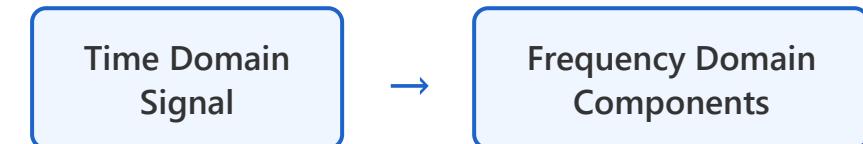
Deep Learning Era: These algorithms are still useful for limited data or lightweight systems.

Audio Feature Extraction: FFT & Spectrogram



Fast Fourier Transform (FFT)

Frequency Domain Transformation



Reveals frequency components present in audio signal



Spectrogram

Visual representation of frequency spectrum over time

Temporal Info

Time

Spectral Info

Frequency

STFT (Short-Time Fourier Transform): Applies FFT to windowed segments



Essential Applications



Audio Analysis



Speech Recognition



Music Processing



Trade-off

Between time and frequency resolution (determined by window size)

Audio Features: MFCC & Chroma



MFCC

Mel-Frequency Cepstral Coefficients

Models human auditory perception

Based on: Mel-scale (perceptually motivated frequency scale)

- Speech recognition
- Speaker identification
- Audio classification



Chroma Features

Pitch Content & Harmonic Structure

Represent pitch content and harmonic structure

Focus on: 12 pitch classes (chromatic scale)

- Music analysis
- Chord recognition
- Similarity detection



Both features capture complementary aspects of audio signals



Compact representations suitable for machine learning models



Visual Process Flow

MFCC Extraction Process

1 Audio Signal

Input original audio signal (time domain)

2 STFT

Frequency analysis with Short-Time Fourier Transform

3 Mel Filter Bank

Apply Mel-scale filters that reflect human auditory characteristics

4 DCT

Extract coefficients with Discrete Cosine Transform

MFCC Coefficient Example (13 coefficients)



Chroma Extraction Process

1 Audio Signal

Input original music signal

2 STFT

Analyze frequency spectrum

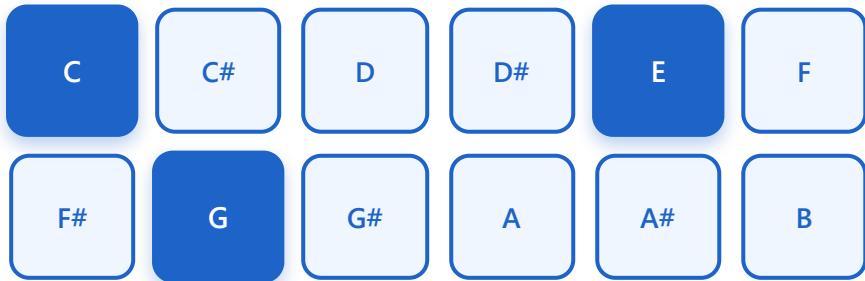
3 Pitch Mapping

Map frequencies to 12 pitch classes (C, C#, D, ...)

4 Octave Folding

Fold all octaves into one (octave invariance)

Chroma Vector Example (12 pitch classes)



Each coefficient represents different frequency characteristics of the audio

Highlighted notes = C Major chord (C, E, G)

The intensity of each pitch class is represented as a vector value

Time Series Statistical Features



Basic Statistics

Mean

Median

Std Dev

Basic statistical properties of the signal



Amplitude Variations

Min

Max

Range

Capture the extent of signal variations



Autocorrelation

Measures signal's correlation with itself at different time lags



Spectral Features

Dominant Frequency

Spectral Entropy

Frequency domain characteristics



Temporal Patterns

Trend

Seasonality

Cyclic Components



Rolling Window Statistics

- Moving average
- Exponential smoothing



Domain-Specific Features

- Peak detection
- Zero-crossing rate

Part 3/3:

Learning-based Representations

- 17.** Concept of Representation Learning
- 18.** Word Embeddings (Word2Vec, GloVe)
- 19.** CNN-based Image Features
- 20.** RNN-based Sequence Features
- 21.** Autoencoders and Latent Representations
- 22.** Transfer Learning Strategies
- 23.** Fine-tuning vs Feature Extraction
- 24.** Domain Adaptation
- 25.** Multimodal Fusion

Concept of Representation Learning

Automatically learns features from raw data instead of manual engineering

Hierarchical Representations

Lower Layers

Capture simple patterns

Higher Layers

Capture complex concepts



End-to-end learning: features optimized for specific task



Reduces need for domain expertise in feature design



Can discover non-obvious patterns humans might miss



Discovers multiple levels of abstraction automatically



Foundation of Modern Deep Learning Success

Word Embeddings: Word2Vec & GloVe

Dense vector representations that capture semantic relationships



Word2Vec

Learns embeddings from local context

CBOW

Skip-gram



GloVe

Learns from global word co-occurrence statistics

Key Properties



Similar words have similar vectors



Much lower dimensions than vocabulary



Transfer learning to downstream tasks



Captures semantic analogies



Classic Analogy Example

king - man + woman ≈ queen



Typical Dimensions: **100-300** (vs. thousands in vocabulary size)

Word2Vec Training Process (Window Size = 2)

CBOW (Continuous Bag of Words)

"The **quick** brown **fox** jumps over "

Step 1: One-Hot Encoding ↓

Context Words (Input):

```
quick: [0,1,0,0,0,...] (vocab size)  
brown: [0,0,1,0,0,...]  
jumps: [0,0,0,1,0,...]  
over: [0,0,0,0,1,...]
```

Input Layer (One-Hot)



Hidden Layer (Embeddings)



Output Layer (Softmax)

Target (Output):

```
fox: [0,0,0,0,0,1,0,...] (one-hot)
```

 Predict center word from averaged context embeddings

Skip-gram

"The quick brown **fox** jumps over"

Step 1: One-Hot Encoding ↓

Center Word (Input):

```
fox: [0,0,0,0,0,1,0,...] (vocab size)
```

Input Layer (One-Hot)



Hidden Layer (Embeddings)



Output Layer (Softmax) ×4

Target (4 Context Words):

```
quick: [0,1,0,0,0,...]  
brown: [0,0,1,0,0,...]  
jumps: [0,0,0,1,0,...]  
over: [0,0,0,0,1,...]
```

 Predict each context word from center word embedding

CNN-based Image Features

Convolutional Neural Networks learn hierarchical visual features

Early Layers

Edges
Textures
Colors

Feature Hierarchy Across Layers

Middle Layers

Object Parts
Patterns

Deep Layers

High-level Semantic
Concepts & Objects



Convolutional filters automatically learned from data



Pooling operations provide translation invariance



Pre-trained CNNs as Feature Extractors

VGG, ResNet, and more serve as powerful feature extractors

Historical Context of CNNs

The evolution of Convolutional Neural Networks in computer vision

1989

LeNet (Yann LeCun)

First successful CNN, used for handwritten digit recognition. Applied to automatic postal code recognition systems

2012

AlexNet (ImageNet Challenge)

Marked the resurgence of deep learning. Achieved overwhelming performance through large-scale GPU training

2014

VGGNet & GoogLeNet

Explored deeper network architectures. VGG featured simple, uniform structure; GoogLeNet introduced Inception modules

2015

ResNet (Residual Networks)

Enabled training of very deep networks (152 layers) with skip connections. Achieved human-level image recognition performance

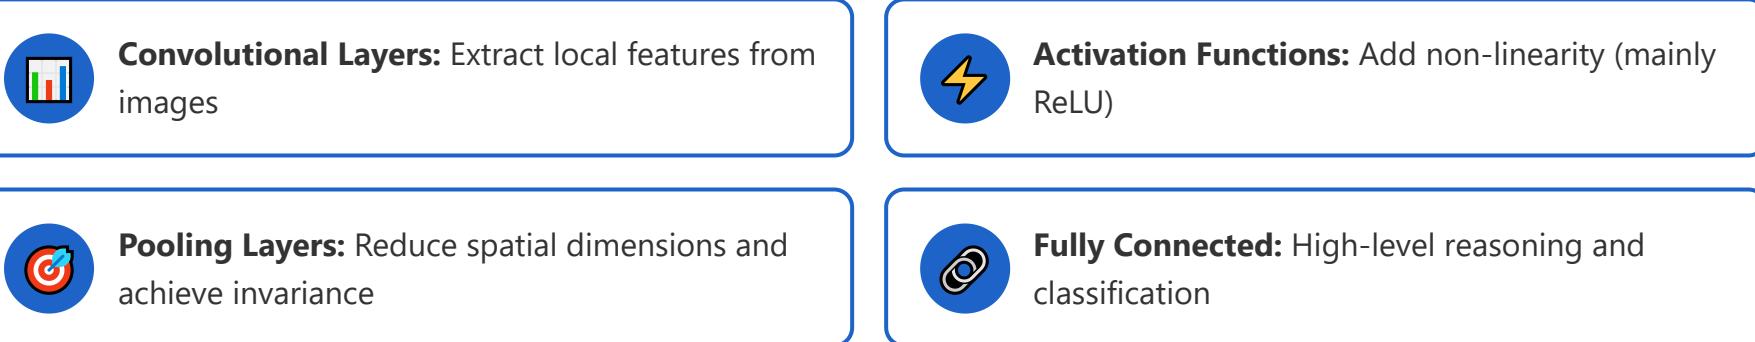
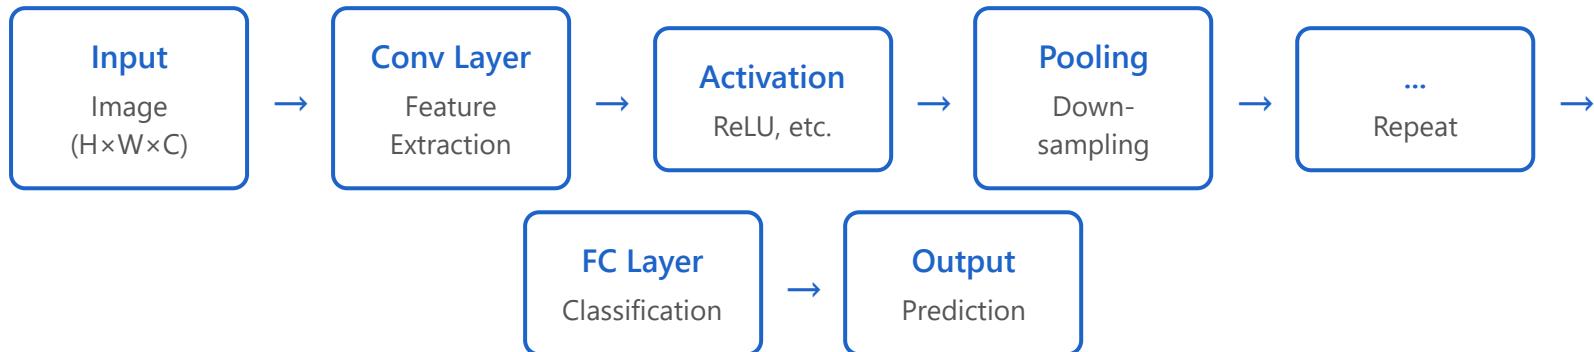
2017+

Modern Architectures

Emergence of modern architectures like EfficientNet and Vision Transformer (ViT) pursuing both efficiency and performance

CNN Layer Architecture

Typical structure of a Convolutional Neural Network



Key Principle: Learn hierarchical features by repeating Conv-Activation-Pooling blocks multiple times

Convolution & Pooling Operations



Convolution Operation

How it Works

A small filter (kernel) slides over the image, performing element-wise multiplication and summation

Key Parameters

- **Filter size:** 3×3 , 5×5 , etc.
- **Stride:** Step size
- **Padding:** Border handling

Features

- Detects local patterns
- Parameter efficiency through weight sharing
- Preserves spatial relationships

Learning

Filter weights are automatically learned from data through backpropagation



Pooling Operation

Purpose

Reduces spatial dimensions of feature maps, decreases computation, and achieves invariance

Max Pooling

Selects maximum value within window. Preserves strongest features. Most widely used

Average Pooling

Computes average value within window. Extracts smoother features

Advantages

- Invariance to small translations/deformations
- Reduces overfitting
- Improves computational efficiency
- Expands receptive field



Together: Convolution extracts features, and Pooling compresses and makes them more robust

Key Concepts in CNN Operations

Receptive Field

- The region of input image that a single neuron "sees"
- Deeper layers have wider receptive fields
- Enables capturing larger contextual information

Output Size Calculation

Output size = $\lfloor (\text{Input size} - \text{Filter size} + 2 \times \text{Padding}) / \text{Stride} \rfloor + 1$

Example: Input 32×32, Filter 5×5, stride 1, padding 2

→ Output: $\lfloor (32 - 5 + 4) / 1 \rfloor + 1 = 32$ (size preserved)

Why CNNs Work Well

1. Translation Invariance:

Can recognize objects even when their position changes

2. Hierarchical Learning:

Simple features → Complex features

3. Parameter Sharing:

Same filter applied across entire image

4. Sparse Connectivity:

Efficient learning through local connections

Feature Visualization

Early layers:

Basic elements like edges, corners, color blobs

Middle layers:

Textures, repetitive patterns, object parts

Deep layers:

Complete objects, scenes, semantic concepts

Filters in each layer learn increasingly complex and abstract features

Interactive Learning Resource

Experience how CNNs work and visually understand the impact of hyperparameters



CNN Explainer

Interactive visualization of Convolutional Neural Networks



[Visit CNN Explainer](#)



Interactive Demo: Visualize what happens in each CNN layer in real-time



Hyperparameters: Directly see the effects of kernel size, stride, padding, etc.



Feature Maps: Observe outputs and activations of each layer in real-time



Educational: The best tool to intuitively understand how CNNs work

Learning Tips

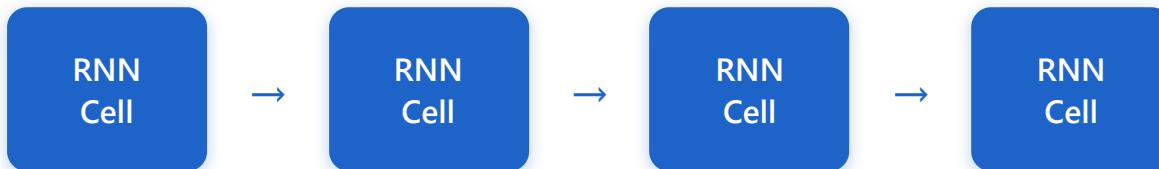
- Start with the "Understanding Hyperparameters" section to learn basic concepts
- Change various hyperparameter values and observe how the output changes

- Follow step-by-step how CNNs respond to real images

RNN-based Sequence Features

Recurrent Neural Networks process sequential data with temporal dependencies

Sequential Processing Architecture



🧠 Hidden states maintain memory of previous inputs

RNN Variants

LSTM

Long Short-Term Memory

GRU

Gated Recurrent Unit

Bidirectional

Past + Future Context

✓ Address vanishing gradient problem



Model variable-length sequences naturally

Applications

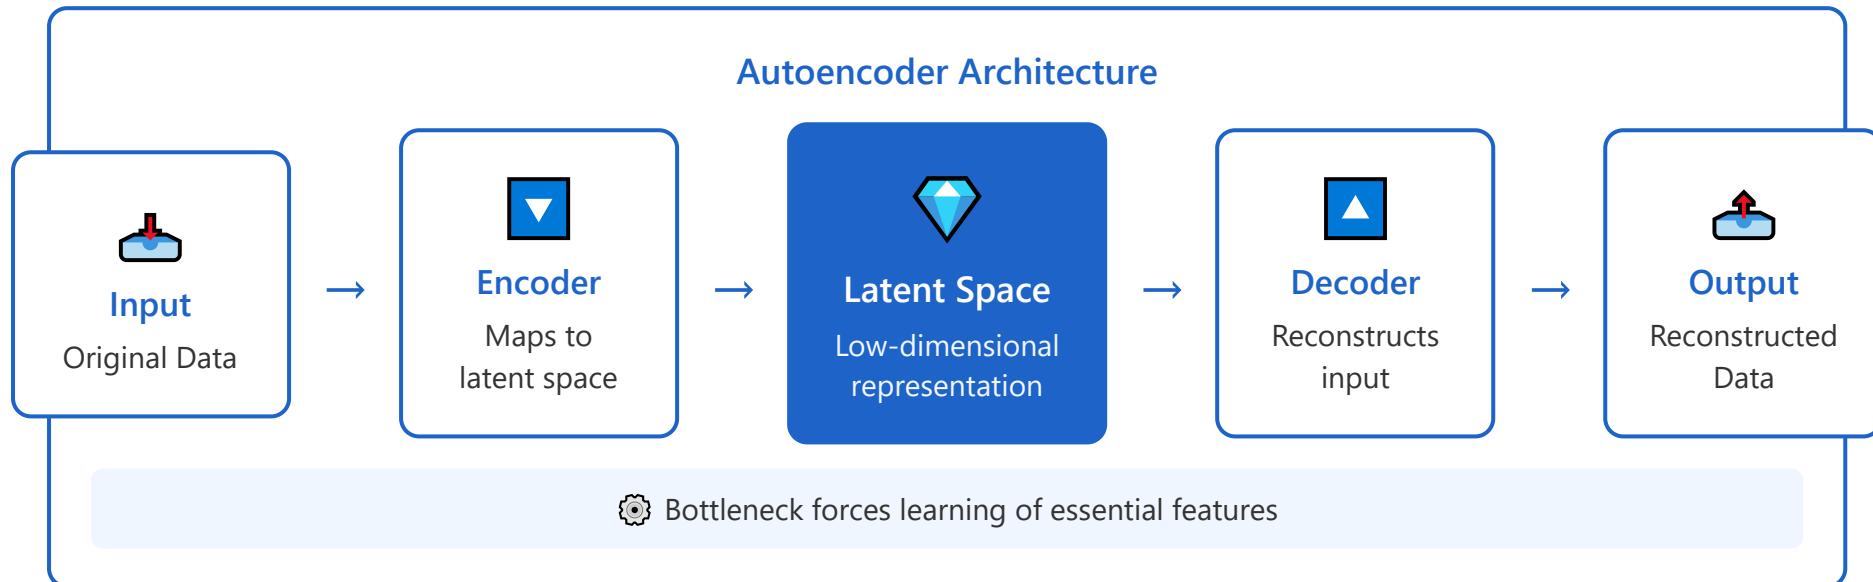
💬 Language Modeling

🗣 Speech Recognition

📈 Time Series

Autoencoders and Latent Representations

Unsupervised learning of compressed data representations



📖 Historical Context

1980s: Hinton and Rumelhart pioneered dimensionality reduction using backpropagation

2006: Hinton's Deep Belief Networks marked the dawn of the deep learning era

2013: Kingma and Welling proposed the Variational Autoencoder (VAE)

Present: Evolved into a core technology for generative models, anomaly detection, and representation learning

🏗 Structural Details

Encoder: $x \rightarrow z = f_{enc}(x; \theta_{enc})$

- Compresses input dimension n to latent dimension d ($d \ll n$)
- Composed of multiple neural network layers (e.g., FC layers, CNNs)

Decoder: $z \rightarrow \hat{x} = f_{\text{dec}}(z; \theta_{\text{dec}})$

- Reconstructs original input from latent representation z
- Can be symmetric or independent from encoder architecture

Loss Function: $L = \|x - \hat{x}\|^2$ (reconstruction error)

Vector Operation Example

Simple Example: 784-dimensional MNIST image \rightarrow 2-dimensional latent space

Input vector: $x \in \mathbb{R}^{784}$ (28×28 image flattened)

e.g., $x = [0.1, 0.0, 0.8, \dots, 0.3]$

Encoder operations:

$h_1 = \text{ReLU}(W_1 x + b_1)$ # $W_1 \in \mathbb{R}^{(128 \times 784)}$

$h_2 = \text{ReLU}(W_2 h_1 + b_2)$ # $W_2 \in \mathbb{R}^{(64 \times 128)}$

$z = W_3 h_2 + b_3$ # $W_3 \in \mathbb{R}^{(2 \times 64)}$, $z \in \mathbb{R}^2$

$\rightarrow z = [1.2, -0.5]$ # 2D latent vector

Decoder operations:

$h_3 = \text{ReLU}(W_4 z + b_4)$ # $W_4 \in \mathbb{R}^{(64 \times 2)}$

$h_4 = \text{ReLU}(W_5 h_3 + b_5)$ # $W_5 \in \mathbb{R}^{(128 \times 64)}$

$\hat{x} = \sigma(W_6 h_4 + b_6)$ # $W_6 \in \mathbb{R}^{(784 \times 128)}$

$\rightarrow \hat{x} = [0.09, 0.02, 0.79, \dots, 0.31]$

Loss: $\text{MSE} = (1/784) \sum (x_i - \hat{x}_i)^2 = 0.003$

Variational Autoencoders (VAE)

Learn probabilistic latent distributions for generation and sampling

- Models latent space as probability distribution in the form $z \sim N(\mu, \sigma^2)$
- Learns regularized latent space with additional KL divergence



Dimensionality
Reduction



Denoising

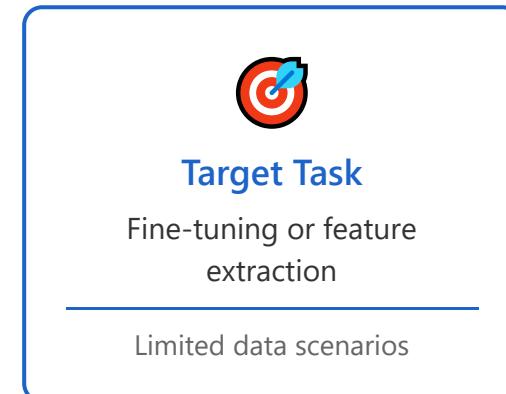
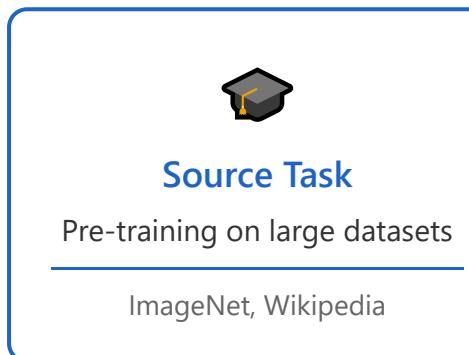


Generation

Transfer Learning Strategies

Leverage knowledge from source task to improve target task performance

Transfer Learning Process



Particularly effective when target task has limited data

Source and target tasks should be related for best results

Reduces training time significantly

Reduces computational requirements



Key Technique in Modern Deep Learning Applications



Transfer Learning Performance Comparison

Training from Scratch

Small Dataset (500 images)

Training Time: 8 hours

Accuracy: 65%

Epochs: 200

VS

Transfer Learning

Same Dataset (500 images)

Training Time: 1 hour

Accuracy: 92%

Epochs: 20



8x faster training and 27% higher accuracy with pre-trained ImageNet model

Fine-tuning vs Feature Extraction



Feature Extraction

Freeze pre-trained model, train only new classifier

- ✓ Faster training
- ✓ Requires less data
- ✓ Lower computational cost
- ✓ Fixed feature representation



Fine-tuning

Update pre-trained weights with small learning rate

- ✓ More task-specific adaptation
- ✓ Better final performance
- ✓ Requires more resources
- ✓ Adapts feature representation



Decision Factors

Dataset Size

Task Similarity

Available Resources



Typical Strategy

1. Feature Extraction



2. Evaluate



3. Fine-tune if needed



Advanced Technique

Layer-wise fine-tuning: Gradually unfreeze deeper layers for progressive adaptation

Domain Adaptation

Addresses distribution shift between training (source) and test (target) domains

Domain Distribution Shift



Source Domain

Training data distribution

Labeled data available



Distribution Shift



Target Domain

Test data distribution

Limited/no labels



Example Scenario

Train on synthetic data → Test on real data



Adaptation Techniques

Domain Adversarial Training

Self-Training

Pseudo-Labeling



Learn domain-invariant representations



Important for real-world deployment

Multimodal Fusion

Combines representations from multiple modalities for unified understanding



Early Fusion

Combine raw features before processing



Intermediate Fusion

Combine learned representations at hidden layers



Late Fusion

Combine decisions from separate modality models

Advanced Techniques



Attention Mechanisms: Focus on relevant modalities



Cross-modal Learning: Use one modality to improve another



Applications

Video Understanding

Visual QA

Multimodal Sentiment Analysis



Improves robustness and performance over single-modality approaches

Thank you

Ho-min Park

homin.park@ghent.ac.kr

powersimmani@gmail.com