

Lecture 19:

Model Explainability - XAI

Fundamentals and Traditional

Methods

Ho-min Park

homin.park@ghent.ac.kr

powersimmani@gmail.com

Lecture Contents

Part 1: Introduction to XAI and its Importance

Part 2: Intrinsically Interpretable Models


Part 3: Feature Importance Methodologies


Part 4: Model-Agnostic Methods


Part 1/4:

Introduction to XAI and Its Importance

- 1.** Course Introduction - The Era of Explainable AI
- 2.** Model Complexity vs Interpretability Trade-off
- 3.** Core Concepts and Terminology in XAI
- 4.** Why Explainability is Necessary
- 5.** XAI Classification Framework
- 6.** XAI Applications by Industry
- 7.** XAI Evaluation Criteria







Why Explainability is Necessary



Trust

Users need to understand AI decisions before accepting them



Accountability

Identifying responsibility when AI makes errors



Debugging

Finding and fixing model biases or data issues



Regulatory Compliance

Legal requirements for automated decision systems



Fairness

Detecting and mitigating discrimination in AI systems



Scientific Discovery


Learning new insights from model patterns



Safety

Ensuring AI behaves as intended in critical applications

XAI Methods



Transparent (White-box)

Black-box (Opaque)

XAI Applications by Industry



Healthcare

Diagnosis predictions, treatment recommendations



Finance

Credit scoring, fraud detection reasoning



Autonomous Vehicles

Decision transparency for safety certification



Criminal Justice

Risk assessment for parole decisions



Human Resources

Fair hiring practices, bias detection




E-commerce

Recommendation system transparency



Manufacturing

XAI Evaluation Dimensions



1 Fidelity

How accurately explanations reflect actual model behavior

2 Comprehensibility

Human ability to understand explanations

3 Consistency

Similar instances receive similar explanations

4 Contrastivity

Explaining why this prediction vs alternatives

5 Actionability

Ability to use explanations for decisions

6 Stability

Robustness to small input perturbations

7 Efficiency

Computational cost of generating explanations

Part 2/4:

Intrinsically Interpretable Models

- 8.** Interpreting Linear Models
- 9.** Transparency of Decision Trees
- 10.** Generalized Additive Models (GAM)
- 11.** Rule-Based Models
- 12.** Monotonic Constraint Models
- 13.** Sparse Linear Models
- 14.** Hands-on: Building Interpretable Models with scikit-learn

Linear Model Interpretation

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n$$

β_0

Intercept

β_i

Coefficients

x_i

Features

Coefficient Magnitude

Indicates feature importance

Sign Direction

Shows positive/negative relationship

Regularization

L1 (Lasso) for selection, L2 (Ridge) for stability

Feature Coefficients

Age

+0.82

Income

+0.63

Education

+0.41

Distance

-0.55


Risk Score

-0.76

Positive Impact

Negative Impact

Decision Tree Structure



If-Then Rules

Binary splits create human-readable decision paths

Complete Path

Full decision logic visible from root to leaf

Feature Importance

Determined by split frequency and information gain

Natural Interactions

Hierarchical splits handle feature interactions

Trade-offs

Deep trees lose interpretability; pruning balances accuracy


- Decision Node
- Approve Outcome
- Reject Outcome

Generalized Additive Models (GAM)

$$y = \beta_0 + f_1(x_1) + f_2(x_2) + f_3(x_3) + \dots$$

where f_i are smooth shape functions

$f_1(\text{Age})$



$f_2(\text{Income})$




$f_3(\text{Credit})$




$f_4(\text{Experience})$



$f_5(\text{Education})$



$f_6(\text{Debt Ratio})$



Key Features

Independent

Each feature contributes separately

Non-linear

Capture complex patterns via smooth functions

Interpretable

Visualize individual feature effects

Separable

Effects remain isolated and analyzable

Balanced

Between linear and black-box models

Rule-Based Decision System

1 RULE

IF Age > 30 **AND** Income > 50K

THEN Decision = APPROVE

Coverage:  85% Confidence:  92%

2 RULE

IF Credit Score > 700 **AND** Debt Ratio < 0.3

THEN Decision = APPROVE

Coverage:  72% Confidence:  88%

3 RULE

IF Age < 25 **AND** Employment < 1 year

Key Concepts

IF-THEN Structure

Rules extracted from data patterns with logical conditions

Coverage

Number of instances satisfying rule conditions

Confidence

Prediction accuracy for instances matching rule

Interpretability

Each prediction explained by triggering rules

Trade-offs

Balancing concise rule sets with high coverage and accuracy

Metrics Guide

THEN Decision = **REJECT**


- █ **Coverage:** % of data covered
- █ **Confidence:** % prediction accuracy

Monotonic Constraint Models

Enforcing Logical Relationships in Predictions



Non-Monotonic Model



✗ Counter-intuitive: Higher score



Monotonic Model



✓ Logical: Higher credit score

Key Features

Logical Constraints

Enforces domain knowledge relationships

Trust Building

Predictions follow expected behavior patterns

Implementation

XGBoost, LightGBM, monotonic NNs

Trade-offs

May slightly reduce accuracy for interpretability

Critical For:

- Regulated industries
- Financial decisions

sometimes decreases approval

consistently increases approval

- Healthcare applications
- High-stakes predictions

Sparse Linear Models

L1 Regularization for Feature Selection

Dense Model

Standard Linear Regression

Age	0.72
Income	0.58
Credit	0.51
Debt	-0.48
Educ.	0.32
Exp.	0.23
Region	0.15
Gender	-0.09

Non-zero
8

Features
8/8

Sparse Model

Lasso (L1 Regularization)

Age	0.68
Income	0.61
Credit	0.54
Debt	-0.45
Educ.	—
Exp.	—
Region	—
Gender	—

Non-zero
4

Features
4/8

Benefits

Auto Selection

Only relevant features retained

Interpretability

Reduced feature count

Implementation

Easier to deploy and debug

High-Dim

Valuable when $p >> n$

Regularization

L1 (Lasso): Sparsity

L2 (Ridge): Stability

Elastic Net: Balanced

Building Interpretable Models with scikit-learn

Hands-on Workflow & Code Examples

Implementation Steps

1 Import & Load Data

Load dataset and dependencies



2 Preprocessing

Feature scaling, categorical encoding



3 Train Models

Linear, Tree-based classifiers

Code Examples

Linear Model

```
from sklearn.linear_model import LogisticRegression  
  
model = LogisticRegression()  
model.fit(X_train, y_train)  
  
# Access coefficients  
coef = model.coef_  
print(f"Coefficients: {coef}")
```

4

Extract Insights

Coefficients, feature importance

**5**

Visualize & Compare

Accuracy vs interpretability



Decision Tree

```
from sklearn.tree import DecisionTreeClassifier  
  
tree = DecisionTreeClassifier(max_depth=3)  
tree.fit(X_train, y_train)  
  
# Feature importance  
importance = tree.feature_importances_
```



Inspection Tools

```
from sklearn.inspection import PartialDependenceDisplay  
  
# Visualize partial dependence  
PartialDependenceDisplay.from_estimator(  
    model, X, features=[0, 1]  
)
```

Key sklearn Modules

linear_model

tree

inspection

preprocessing

Part 3/4:


Feature Importance Methodologies

- 15.** Permutation Importance
- 16.** Drop-Column Importance
- 17.** Partial Dependence Plots (PDP)
- 18.** Individual Conditional Expectation (ICE)
- 19.** Accumulated Local Effects (ALE)
- 20.** Feature Interaction Analysis

Permutation Importance

Feature Importance via Random Shuffling

Top 10 Features by Importance Score



How It Works

Process

- 1 Measure baseline performance
- 2 Shuffle feature values randomly
- 3 Re-evaluate model performance
- 4 Importance = Performance drop

Model-Agnostic

Works with any trained model

Interactions

Captures feature contributions including interactions

Implementation

`sklearn.inspection.permutation_importance()`

✓ Advantages

Simple, reliable, no assumptions

⚠ Limitations

Correlated features may deflate

Permutation Importance: Detailed Algorithm (Example)

1 Calculate Baseline Score

Evaluate the trained model on validation dataset to establish baseline performance metric

2 Select Feature to Test

Choose one feature column from the dataset to permute

```
baseline_score = model.score(X_val, y_val) →  
baseline_score = 0.850 (R2 score)
```

```
feature_to_test = 'Credit Score' # Testing most important  
feature first
```

3 Randomly Shuffle Feature

Randomly permute the values of selected feature, breaking its relationship with target

```
X_permuted = X_val.copy() X_permuted['Credit Score'] =  
shuffle([750, 680, 820, ...]) → [820, 750, 680, ...] #  
Randomly shuffled
```

4 Re-evaluate Performance

Calculate new performance score with shuffled feature values

```
permuted_score = model.score(X_permuted, y_val) →  
permuted_score = 0.765 (R2 dropped!)
```

5 Calculate Importance

Feature importance = Performance drop caused by shuffling

```
importance = baseline_score - permuted_score  
importance =  
0.850 - 0.765 = 0.085 ✓
```

6 Repeat for Stability

Repeat shuffling multiple times and average for stable estimate

```
importances = [0.085, 0.082, 0.088, 0.084, 0.086]  
mean_importance = 0.085 ± 0.002
```

7 Iterate All Features

Repeat process for every feature to get complete ranking

```
Credit Score: 0.085 Income: 0.072 Age: 0.065 Debt Ratio:  
0.058 ...
```

8 Rank & Visualize

Sort features by importance and create visualization with error bars

```
sorted_features = sort_by_importance() → See bar chart  
above! 📈
```

Drop-Column Importance

Feature Removal Impact Analysis

Importance Score Comparison

Feature	Permutation	Drop-Column	Δ Diff
Credit Score	0.085	0.092	+0.007
Income	0.072	0.078	+0.006
Age	0.065	0.048	-0.017
Debt Ratio	0.058	0.061	+0.003
Employment	0.048	0.032	-0.016
Education	0.042	0.044	+0.002
Experience	0.035	0.036	+0.001
Savings	0.028	0.028	0.000

Method

Process

- 1 Train with all features
- 2 Remove one feature
- 3 Retrain model
- 4 Measure performance drop

Accuracy

More accurate than permutation

Cost

Requires n retrainings for n features

Correlations

Handles feature correlations better

Use Case

Final feature selection decisions


Key Insight

Fast

Accurate

Highlighted rows show redundant features

Calculation Example




Interpretation: Removing "Age" causes a 4.8% drop in accuracy, indicating it's a moderately important feature. Compare this across all features to identify which ones are truly essential for your model. Features showing lower importance in Drop-Column than Permutation (like Age: -0.017) might be redundant when other correlated features are present.


Partial Dependence Plots (PDP)

Marginal Effect Visualization

1D PDP: Age vs Loan Approval Probability



2D PDP: Age × Income Interaction



Key Concepts

Marginal Effect

Average prediction as feature varies

1D PDP

Single feature effect (line plot)

2D PDP

Two-feature interactions (heatmap)

Implementation

`sklearn.inspection.PartialDependenceDisplay`

⚠ Assumption

Features are independent - misleading if violated

Reveals

- Non-linear relationships
- Thresholds & transitions
- Average behavior patterns

How PDP is Calculated

Step-by-Step Mathematical Process

$$PDP_{xs}(x_s) = E_{X_c}[f(x_s, X_c)] = (1/n) \sum_{i=1}^n f(x_s, x_c^{(i)})$$

x_s : Target feature value (fixed) | X_c : All other features (vary across data) | f : Model prediction function | n : Number of data points

1

Original Data

Age	Inc	\hat{y}
25	40K	0.3
35	60K	0.7
45	80K	0.9
55	50K	0.6

Start with training dataset containing all features and their predictions

2

Fix Target Feature

Age	Inc	\hat{y}
30	40K	?
30	60K	?
30	80K	?
30	50K	?

Replace target feature (Age) with specific value (e.g., 30) for ALL samples

3

Predict & Average


$$\begin{aligned} f(30, 40K) &= 0.45 \\ f(30, 60K) &= 0.65 \\ f(30, 80K) &= 0.75 \\ f(30, 50K) &= 0.55 \end{aligned}$$

$$\begin{aligned} PDP(\text{Age}=30) &= \\ (0.45+0.65+0.75+0.55)/4 &= 0.60 \end{aligned}$$

Get predictions for all modified samples, then compute their average

4

Repeat & Plot



Repeat steps 2-3 for different Age values to create the complete PDP curve


Key Insight

PDP marginalizes out other features by averaging predictions across their distribution. This shows the **average effect** of the target feature, assuming feature independence.

Individual Conditional Expectation (ICE)

Instance-Level Feature Effects

ICE Curves: Age vs Loan Approval Prediction



Key Concepts

ICE vs PDP

ICE

Individual instance trajectories

PDP

Average of all ICE curves

Heterogeneity

Reveals variation in feature effects across instances

Subgroups

Identifies different feature-outcome relationships

Interactions

Detects interactions PDP may hide

Non-parallel Lines

Shows unexpected patterns & heterogeneity

Centered ICE

Subtract baseline for better comparison

What to Look For

- Parallel lines = homogeneous effect
- Diverging lines = heterogeneous effect

How ICE Works: Step-by-Step Process

1

Select Instance

Choose a single data point from your dataset

```
x1 = (age=30,  
income=50k, ...)
```

2

Vary Feature

Change target feature across its range, keep others fixed

```
age: 20, 25, 30,  
..., 65  
income=50k (fixed)
```

3

Get Predictions

Run model for each feature value

```
ŷ(20), ŷ(25), ŷ(30),  
..., ŷ(65)
```

4

Plot & Repeat

Draw curve for this instance, repeat for all instances

```
n curves for n  
instances  
Average = PDP
```

Accumulated Local Effects (ALE)

Unbiased Feature Effects with Correlated Features



PDP (Biased)

✓ ALE (Unbiased)

Key Features

ALE Advantages

- ✓ Unbiased with correlated features
- ✓ Local neighborhoods only
- ✓ Computationally efficient
- ✓ No marginalization needed

Local Effects

Accumulated across feature range

Reliability

More trustworthy for realistic datasets


Implementation

ALEPython package

Quick Comparison

	PDP	ALE
Correlated	Biased	Unbiased
Speed	Slow	Fast


Effect



How

intervals. E
ally exists.

Effect



Differences

predictions when the feature value
interval. This represents the "local

$\}) = f(z_{-j}, X_{\{-j\}})$
nts in each interval

Average the local effects of all data points within each interval. This cancels out the effects of other features.

$$ALE_j = (1/n_j) \sum [f(z_{j+1}, X_i) - f(z_j, X_i)]$$

n_j : number of data points in interval j

ys i

Starting from the first interval, accumulate the local effects of each interval. This is what "Accumulated" in Accumulated Local Effects means.

$$ALE(z) = \sum_{j=1}^k ALE_j$$

Cumulative from left to right

5

Center the Plot

Center the entire ALE plot so the average is 0. This allows us to see the relative effect of the feature rather than its absolute effect.

```
ALE_centered(z) = ALE(z) - E[ALE(z)]  
Interpretation: effect relative to average
```

6

Visualize & Interpret

The final ALE plot shows the pure effect on model predictions at each feature value. An upward slope indicates positive effect, while a downward slope indicates negative effect.




Key Takeaway

ALE operates only within **local regions where actual data exists**, so it doesn't suffer from the bias caused by unrealistic data combinations like PDP does. Even when features are correlated, ALE measures pure effects while maintaining the natural distribution of other features by considering only small changes within each interval.

Feature Interaction Analysis

Pairwise Interaction Strength Matrix



Methods

Interaction Detection

- **H-statistic:** Quantify strength
- **2D PDP:** Visual effects
- **SHAP:** Pairwise values
- **PD-based:** Joint vs separate

Non-additive

Features work together synergistically

Understanding

Model behavior & debugging

Engineering

Guide interaction term creation

Complexity

$O(n^2)$ - quadratic with features

Interpretation Guide

- 0.0-0.3: Weak interaction
- 0.3-0.6: Moderate interaction
- 0.6-1.0: Strong interaction


Part 4/4:

Model-Agnostic Methods

- 21.** Surrogate Models
- 22.** Introduction to LIME
- 23.** LIME Advanced Topics
- 24.** Anchor Explanations
- 25.** Practical Guidelines and Best Practices

Surrogate Models

Approximating Black-Box with Interpretable Models



Interpret Surrogate

Extract insights and explanations

- Decision Trees
- Linear Models
- Rule Sets
- Generalized Additive Models

Introduction to LIME

Local Interpretable Model-agnostic Explanations


LIME Process (4 Steps)

1 Original Instance



Instance to explain

2 Perturbed Samples



Generate neighbors

3 Local Linear Model



4 Feature Importance



Key Ideas

Core Concepts

- **Local:** Explains one prediction
- **Proximity:** Weight by distance
- **Linear:** Interpretable locally
- **Agnostic:** Any model type

Model-Agnostic

Works with any ML model

Perturbation

Create samples around instance

Weighting

Closer samples matter more

Interpretability

Feature importance shows factors

Python Implementation

Fit weighted model

Extract explanation

lime package:





- lime.lime_tabular
- lime.lime_text
- lime.lime_image

LIME Advanced Topics

Stability, Hyperparameters, and Limitations

Stability Issue Demonstration

⚠ Same instance, different random seeds → different explanations



Key Topics

Challenges

- ⚠ Stability issues
- ⚠ Inconsistent explanations
- ⚠ Sampling sensitivity

Hyperparameters

Kernel width, sample size

Feature Selection

Subset for simplicity

Categorical Vars

Special perturbation strategies

SP-LIME

Representative instances

Text & Image

Segment-based perturbations

Best Practices

- ✓ Run multiple times

Credit +0.20

Debt

-0.23



✓ Compare with other methods

Anchor Explanations

Sufficient Conditions with High Precision

ANCHOR RULE EXAMPLE

IF Age > 30 **AND** Income > 50K **THEN** Approve



Key Features

Advantages

- ✓ More stable than LIME
- ✓ Human-readable rules
- ✓ High confidence regions
- ✓ Less perturbation-sensitive

Sufficient Conditions

Rules that guarantee prediction

Precision Metric

Rule accuracy on samples

Coverage Metric


Fraction where rule applies

Beam Search

Finds optimal anchor rules

Implementation


Python: anchor package
(integrated with Alibi library)



Practical Guidelines and Best Practices

Selecting the Right XAI Method

XAI Method Selection Decision Tree



Best Practices

Key Guidelines

- Use multiple XAI methods
- Validate with domain experts
- Consider computational budget
- Document limitations
- Match stakeholder needs

Validation

Match domain expertise

Budget

Real-time vs batch processing

Documentation


Record assumptions & limits

Monitoring

Track explanation drift

⚠ Watch Out For

- ! Adversarial examples
- ! Data distribution shifts
- ! Over-reliance on single method



Thank you

Ho-min Park

homin.park@ghent.ac.kr

powersimmani@gmail.com