

Lecture 19:

# **Model Explainability - XAI**

## **Fundamentals and Traditional**

## **Methods**

**Ho-min Park**

[homin.park@ghent.ac.kr](mailto:homin.park@ghent.ac.kr)

[powersimmani@gmail.com](mailto:powersimmani@gmail.com)

# Lecture Contents

**Part 1:** Introduction to XAI and its Importance

**Part 2:** Intrinsically Interpretable Models


**Part 3:** Feature Importance Methodologies


**Part 4:** Model-Agnostic Methods


**Part 1/4:**

# **Introduction to XAI and Its Importance**

- 1.** Course Introduction - The Era of Explainable AI
- 2.** Model Complexity vs Interpretability Trade-off
- 3.** Core Concepts and Terminology in XAI
- 4.** Why Explainability is Necessary
- 5.** XAI Classification Framework
- 6.** XAI Applications by Industry
- 7.** XAI Evaluation Criteria







## Why Explainability is Necessary



### Trust

Users need to understand AI decisions before accepting them



### Accountability

Identifying responsibility when AI makes errors



### Debugging

Finding and fixing model biases or data issues



### Regulatory Compliance

Legal requirements for automated decision systems



### Fairness

Detecting and mitigating discrimination in AI systems



### Scientific Discovery


Learning new insights from model patterns



### Safety

Ensuring AI behaves as intended in critical applications

## XAI Methods



Transparent (White-box)

Black-box (Opaque)

## XAI Applications by Industry



### Healthcare

Diagnosis predictions, treatment recommendations



### Finance

Credit scoring, fraud detection reasoning



### Autonomous Vehicles

Decision transparency for safety certification



### Criminal Justice

Risk assessment for parole decisions



### Human Resources

Fair hiring practices, bias detection



### E-commerce


Recommendation system transparency



### Manufacturing



## XAI Evaluation Dimensions



### 1 Fidelity

How accurately explanations reflect actual model behavior

### 2 Comprehensibility

Human ability to understand explanations

### 3 Consistency

Similar instances receive similar explanations

### 4 Contrastivity

Explaining why this prediction vs alternatives

### 5 Actionability

Ability to use explanations for decisions

### 6 Stability

Robustness to small input perturbations

### 7 Efficiency

Computational cost of generating explanations

## **Part 2/4:**

# **Intrinsically Interpretable Models**

- 8.** Interpreting Linear Models
- 9.** Transparency of Decision Trees
- 10.** Generalized Additive Models (GAM)
- 11.** Rule-Based Models
- 12.** Monotonic Constraint Models
- 13.** Sparse Linear Models
- 14.** Hands-on: Building Interpretable Models with scikit-learn

## Linear Model Interpretation

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n$$

$\beta_0$

Intercept

$\beta_i$

Coefficients

$x_i$

Features

### Coefficient Magnitude

Indicates feature importance

### Sign Direction

Shows positive/negative relationship

### Regularization

L1 (Lasso) for selection, L2 (Ridge) for stability

## Feature Coefficients

Age

+0.82

Income

+0.63

Education

+0.41

Distance

-0.55


Risk Score

-0.76

Positive Impact

Negative Impact

## Decision Tree Structure



### If-Then Rules

Binary splits create human-readable decision paths

### Complete Path

Full decision logic visible from root to leaf

### Feature Importance

Determined by split frequency and information gain

### Natural Interactions

Hierarchical splits handle feature interactions

### Trade-offs

Deep trees lose interpretability; pruning balances accuracy


- Decision Node
- Approve Outcome
- Reject Outcome

## Generalized Additive Models (GAM)

$$y = \beta_0 + f_1(x_1) + f_2(x_2) + f_3(x_3) + \dots$$

where  $f_i$  are smooth shape functions

$f_1(\text{Age})$



$f_2(\text{Income})$




$f_3(\text{Credit})$




$f_4(\text{Experience})$



$f_5(\text{Education})$



$f_6(\text{Debt Ratio})$



### Key Features

#### Independent

Each feature contributes separately

#### Non-linear

Capture complex patterns via smooth functions

#### Interpretable

Visualize individual feature effects

#### Separable

Effects remain isolated and analyzable

#### Balanced

Between linear and black-box models

## Rule-Based Decision System

### 1 RULE

**IF** Age > 30 **AND** Income > 50K

**THEN** Decision = APPROVE

Coverage:  85% Confidence:  92%

### 2 RULE

**IF** Credit Score > 700 **AND** Debt Ratio < 0.3

**THEN** Decision = APPROVE

Coverage:  72% Confidence:  88%

### 3 RULE

**IF** Age < 25 **AND** Employment < 1 year

## Key Concepts

### IF-THEN Structure

Rules extracted from data patterns with logical conditions

### Coverage

Number of instances satisfying rule conditions

### Confidence

Prediction accuracy for instances matching rule

### Interpretability

Each prediction explained by triggering rules

### Trade-offs

Balancing concise rule sets with high coverage and accuracy

## Metrics Guide

**THEN** Decision = **REJECT**

- █ **Coverage:** % of data covered
- █ **Confidence:** % prediction accuracy

## Monotonic Constraint Models

*Enforcing Logical Relationships in Predictions*



### Non-Monotonic Model



✗ Counter-intuitive: Higher score



### Monotonic Model



✓ Logical: Higher credit score

### Key Features

#### Logical Constraints

Enforces domain knowledge relationships

#### Trust Building

Predictions follow expected behavior patterns

#### Implementation

XGBoost, LightGBM, monotonic NNs

#### Trade-offs

May slightly reduce accuracy for interpretability

#### Critical For:

- Regulated industries
- Financial decisions

sometimes decreases approval

consistently increases approval

- Healthcare applications
- High-stakes predictions

## Sparse Linear Models

L1 Regularization for Feature Selection

### Dense Model

Standard Linear Regression

Age	0.72
Income	0.58
Credit	0.51
Debt	-0.48
Educ.	0.32
Exp.	0.23
Region	0.15
Gender	-0.09

Non-zero  
**8**

Features  
**8/8**

### Sparse Model

Lasso (L1 Regularization)

Age	0.68
Income	0.61
Credit	0.54
Debt	-0.45
Educ.	—
Exp.	—
Region	—
Gender	—

Non-zero  
**4**

Features  
**4/8**

### Benefits

#### Auto Selection

Only relevant features retained

#### Interpretability

Reduced feature count

#### Implementation

Easier to deploy and debug

#### High-Dim

Valuable when  $p >> n$

### Regularization

**L1 (Lasso):** Sparsity

**L2 (Ridge):** Stability

**Elastic Net:** Balanced

# Building Interpretable Models with scikit-learn

*Hands-on Workflow & Code Examples*

## Implementation Steps

### 1 Import & Load Data

Load dataset and dependencies



### 2 Preprocessing

Feature scaling, categorical encoding



### 3 Train Models

Linear, Tree-based classifiers

## Code Examples

### Linear Model

```
from sklearn.linear_model import LogisticRegression  
  
model = LogisticRegression()  
model.fit(X_train, y_train)  
  
# Access coefficients  
coef = model.coef_  
print(f"Coefficients: {coef}")
```

**4**

## Extract Insights

Coefficients, feature importance

**5**

## Visualize & Compare

Accuracy vs interpretability



### Decision Tree

```
from sklearn.tree import DecisionTreeClassifier  
  
tree = DecisionTreeClassifier(max_depth=3)  
tree.fit(X_train, y_train)  
  
# Feature importance  
importance = tree.feature_importances_
```



### Inspection Tools

```
from sklearn.inspection import PartialDependenceDisplay  
  
# Visualize partial dependence  
PartialDependenceDisplay.from_estimator(  
    model, X, features=[0, 1]  
)
```

### Key sklearn Modules

linear\_model

tree

inspection

preprocessing

**Part 3/4:**


# **Feature Importance Methodologies**

- 15.** Permutation Importance
- 16.** Drop-Column Importance
- 17.** Partial Dependence Plots (PDP)
- 18.** Individual Conditional Expectation (ICE)
- 19.** Accumulated Local Effects (ALE)
- 20.** Feature Interaction Analysis

## Permutation Importance

Feature Importance via Random Shuffling

### Top 10 Features by Importance Score



## How It Works

### Process

- 1 Measure baseline performance
- 2 Shuffle feature values randomly
- 3 Re-evaluate model performance
- 4 Importance = Performance drop

### Model-Agnostic

Works with any trained model

### Interactions

Captures feature contributions including interactions

### Implementation

`sklearn.inspection.permutation_importance()`

### ✓ Advantages

Simple, reliable, no assumptions

### ⚠ Limitations

Correlated features may deflate

## Permutation Importance: Detailed Algorithm (Example)

### 1 Calculate Baseline Score

Evaluate the trained model on validation dataset to establish baseline performance metric

### 2 Select Feature to Test

Choose one feature column from the dataset to permute

```
baseline_score = model.score(X_val, y_val) →  
baseline_score = 0.850 (R2 score)
```

```
feature_to_test = 'Credit Score' # Testing most important  
feature first
```

### 3 Randomly Shuffle Feature

Randomly permute the values of selected feature, breaking its relationship with target

```
X_permuted = X_val.copy() X_permuted['Credit Score'] =  
shuffle([750, 680, 820, ...]) → [820, 750, 680, ...] #  
Randomly shuffled
```

### 4 Re-evaluate Performance

Calculate new performance score with shuffled feature values

```
permuted_score = model.score(X_permuted, y_val) →  
permuted_score = 0.765 (R2 dropped!)
```

### 5 Calculate Importance

Feature importance = Performance drop caused by shuffling

```
importance = baseline_score - permuted_score  
importance =  
0.850 - 0.765 = 0.085 ✓
```

### 6 Repeat for Stability

Repeat shuffling multiple times and average for stable estimate

```
importances = [0.085, 0.082, 0.088, 0.084, 0.086]  
mean_importance = 0.085 ± 0.002
```

### 7 Iterate All Features

Repeat process for every feature to get complete ranking

```
Credit Score: 0.085 Income: 0.072 Age: 0.065 Debt Ratio:  
0.058 ...
```

### 8 Rank & Visualize

Sort features by importance and create visualization with error bars

```
sorted_features = sort_by_importance() → See bar chart  
above! 📈
```

# Drop-Column Importance

Feature Removal Impact Analysis

## Importance Score Comparison

Feature	Permutation	Drop-Column	Δ Diff
Credit Score	0.085	0.092	+0.007
Income	0.072	0.078	+0.006
Age	0.065	0.048	-0.017
Debt Ratio	0.058	0.061	+0.003
Employment	0.048	0.032	-0.016
Education	0.042	0.044	+0.002
Experience	0.035	0.036	+0.001
Savings	0.028	0.028	0.000

## Method

### Process

- 1 Train with all features
- 2 Remove one feature
- 3 Retrain model
- 4 Measure performance drop

### Accuracy

More accurate than permutation

### Cost

Requires n retrainings for n features

### Correlations

Handles feature correlations better

### Use Case

Final feature selection decisions


### Key Insight

Fast

Accurate

Highlighted rows show redundant features

## Calculation Example




**Interpretation:** Removing "Age" causes a 4.8% drop in accuracy, indicating it's a moderately important feature. Compare this across all features to identify which ones are truly essential for your model. Features showing lower importance in Drop-Column than Permutation (like Age: -0.017) might be redundant when other correlated features are present.


# Partial Dependence Plots (PDP)

Marginal Effect Visualization

## 1D PDP: Age vs Loan Approval Probability



## 2D PDP: Age × Income Interaction



## Key Concepts

### Marginal Effect

Average prediction as feature varies

### 1D PDP

Single feature effect (line plot)

### 2D PDP

Two-feature interactions (heatmap)

### Implementation

`sklearn.inspection.PartialDependenceDisplay`

### ⚠ Assumption

Features are independent - misleading if violated

### Reveals

- Non-linear relationships
- Thresholds & transitions
- Average behavior patterns

## How PDP is Calculated

Step-by-Step Mathematical Process

$$PDP_{xs}(x_s) = E_{X_c}[f(x_s, X_c)] = (1/n) \sum_{i=1}^n f(x_s, x_c^{(i)})$$

$x_s$ : Target feature value (fixed) |  $X_c$ : All other features (vary across data) |  $f$ : Model prediction function |  $n$ : Number of data points

1

### Original Data

Age	Inc	$\hat{y}$
25	40K	0.3
35	60K	0.7
45	80K	0.9
55	50K	0.6

Start with training dataset containing all features and their predictions

2

### Fix Target Feature

Age	Inc	$\hat{y}$
30	40K	?
30	60K	?
30	80K	?
30	50K	?

Replace target feature (Age) with specific value (e.g., 30) for ALL samples

3

### Predict & Average


$$\begin{aligned} f(30, 40K) &= 0.45 \\ f(30, 60K) &= 0.65 \\ f(30, 80K) &= 0.75 \\ f(30, 50K) &= 0.55 \end{aligned}$$

$$\begin{aligned} PDP(\text{Age}=30) &= \\ (0.45+0.65+0.75+0.55)/4 &= 0.60 \end{aligned}$$

Get predictions for all modified samples, then compute their average

4

### Repeat & Plot



Repeat steps 2-3 for different Age values to create the complete PDP curve


#### Key Insight

PDP marginalizes out other features by averaging predictions across their distribution. This shows the **average effect** of the target feature, assuming feature independence.

# Individual Conditional Expectation (ICE)

Instance-Level Feature Effects

ICE Curves: Age vs Loan Approval Prediction



## Key Concepts

### ICE vs PDP

#### ICE

Individual instance trajectories

#### PDP

Average of all ICE curves

### Heterogeneity

Reveals variation in feature effects across instances

### Subgroups

Identifies different feature-outcome relationships

### Interactions

Detects interactions PDP may hide

### Non-parallel Lines

Shows unexpected patterns & heterogeneity

### Centered ICE

Subtract baseline for better comparison

### What to Look For

- Parallel lines = homogeneous effect
- Diverging lines = heterogeneous effect

## How ICE Works: Step-by-Step Process

**1**

### Select Instance

Choose a single data point from your dataset

```
x1 = (age=30,  
income=50k, ...)
```

**2**

### Vary Feature

Change target feature across its range, keep others fixed

```
age: 20, 25, 30,  
..., 65  
income=50k (fixed)
```

**3**

### Get Predictions

Run model for each feature value

```
ŷ(20), ŷ(25), ŷ(30),  
..., ŷ(65)
```

**4**

### Plot & Repeat

Draw curve for this instance, repeat for all instances

```
n curves for n  
instances  
Average = PDP
```

# Accumulated Local Effects (ALE)

*Unbiased Feature Effects with Correlated Features*



PDP (Biased)

✓ ALE (Unbiased)

## Key Features

### ALE Advantages

- ✓ Unbiased with correlated features
- ✓ Local neighborhoods only
- ✓ Computationally efficient
- ✓ No marginalization needed

### Local Effects

Accumulated across feature range

### Reliability


More trustworthy for realistic datasets

### Implementation

ALEPython package


### Quick Comparison

	PDP	ALE
Correlated	Biased	Unbiased
Speed	Slow	Fast



## How

intervals. E  
ally exists.



## Differences

predictions when the feature value  
interval. This represents the "local

$\}) = f(z_{-j}, x_{\{-j\}})$   
nts in each interval

Average the local effects of all data points within each interval. This cancels out the effects of other features.

$$ALE_j = (1/n_j) \sum [f(z_{j+1}, x_i) - f(z_j, x_i)]$$

n\_j: number of data points in interval j

ys i

Starting from the first interval, accumulate the local effects of each interval. This is what "Accumulated" in Accumulated Local Effects means.

$$ALE(z) = \sum_{j=1}^k ALE_j$$

Cumulative from left to right

## 5

## Center the Plot

Center the entire ALE plot so the average is 0. This allows us to see the relative effect of the feature rather than its absolute effect.

```
ALE_centered(z) = ALE(z) - E[ALE(z)]  
Interpretation: effect relative to average
```

## 6

## Visualize & Interpret

The final ALE plot shows the pure effect on model predictions at each feature value. An upward slope indicates positive effect, while a downward slope indicates negative effect.




### Key Takeaway

ALE operates only within **local regions where actual data exists**, so it doesn't suffer from the bias caused by unrealistic data combinations like PDP does. Even when features are correlated, ALE measures pure effects while maintaining the natural distribution of other features by considering only small changes within each interval.

# Feature Interaction Analysis

Pairwise Interaction Strength Matrix



## Methods

### Interaction Detection

- **H-statistic:** Quantify strength
- **2D PDP:** Visual effects
- **SHAP:** Pairwise values
- **PD-based:** Joint vs separate

### Non-additive

Features work together synergistically

### Understanding

Model behavior & debugging

### Engineering

Guide interaction term creation

### Complexity

$O(n^2)$  - quadratic with features

### Interpretation Guide

- 0.0-0.3: Weak interaction
- 0.3-0.6: Moderate interaction
- 0.6-1.0: Strong interaction

**Part 4/4:**


# **Model-Agnostic Methods**

- 21.** Surrogate Models
- 22.** Introduction to LIME
- 23.** LIME Advanced Topics
- 24.** Anchor Explanations
- 25.** Practical Guidelines and Best Practices

# Surrogate Models

*Approximating Black-Box with Interpretable Models*

## Surrogate Model Process



## Key Concepts

### Types

#### Global Surrogate

Mimic entire model behavior

#### Local Surrogate

Approximate specific region

### Fidelity

How well surrogate matches original

### Advantages

Leverage simple model interpretability

### Limitations

May not capture full complexity

### Use Cases

Compression, explanation, debugging

### Common Surrogates

## Interpret Surrogate

Extract insights and explanations

- Decision Trees
- Linear Models
- Rule Sets
- Generalized Additive Models

# Introduction to LIME

*Local Interpretable Model-agnostic Explanations*


## LIME Process (4 Steps)

### 1 Original Instance



Instance to explain

### 2 Perturbed Samples



Generate neighbors

### 3 Local Linear Model



### 4 Feature Importance



## Key Ideas

### Core Concepts

- **Local:** Explains one prediction
- **Proximity:** Weight by distance
- **Linear:** Interpretable locally
- **Agnostic:** Any model type

### Model-Agnostic

Works with any ML model

### Perturbation

Create samples around instance

### Weighting

Closer samples matter more

### Interpretability

Feature importance shows factors

### Python Implementation

Fit weighted model

Extract explanation

**lime** package:





- lime.lime\_tabular
- lime.lime\_text
- lime.lime\_image

# LIME Advanced Topics

*Stability, Hyperparameters, and Limitations*

## Stability Issue Demonstration

⚠ Same instance, different random seeds → different explanations



## Key Topics

### Challenges

- ⚠ Stability issues
- ⚠ Inconsistent explanations
- ⚠ Sampling sensitivity

### Hyperparameters

Kernel width, sample size

### Feature Selection

Subset for simplicity

### Categorical Vars

Special perturbation strategies

### SP-LIME

Representative instances

### Text & Image

Segment-based perturbations

### Best Practices

- ✓ Run multiple times

Credit +0.20

Debt

-0.23


✓ Compare with other methods

## Anchor Explanations

*Sufficient Conditions with High Precision*

### ANCHOR RULE EXAMPLE

**IF** Age > 30 **AND** Income > 50K **THEN** Approve



### Key Features

#### Advantages

- ✓ More stable than LIME
- ✓ Human-readable rules
- ✓ High confidence regions
- ✓ Less perturbation-sensitive

#### Sufficient Conditions

Rules that guarantee prediction

#### Precision Metric

Rule accuracy on samples

#### Coverage Metric


Fraction where rule applies

#### Beam Search

Finds optimal anchor rules

#### Implementation


**Python:** anchor package  
(integrated with Alibi library)



# Practical Guidelines and Best Practices

Selecting the Right XAI Method

## XAI Method Selection Decision Tree



## Best Practices

### Key Guidelines

- Use multiple XAI methods
- Validate with domain experts
- Consider computational budget
- Document limitations
- Match stakeholder needs

### Validation

Match domain expertise

### Budget

Real-time vs batch processing

### Documentation


Record assumptions & limits

### Monitoring

Track explanation drift

### ⚠ Watch Out For

- ! Adversarial examples
- ! Data distribution shifts
- ! Over-reliance on single method



# Thank you

Ho-min Park

[homin.park@ghent.ac.kr](mailto:homin.park@ghent.ac.kr)

[powersimmani@gmail.com](mailto:powersimmani@gmail.com)