Batch Size and Memory Usage Calculation

Memory Usage Components

Model Params + Optimizer States + Gradients + Activations

Activations Memory $\propto$ **Batch Size × Sequence Length**

Example: BERT-base

batch = 32 • **seq = 512** → **~8GB** activation memory

Gradient Accumulation

Split batch into micro-batches to save memory

Gradient Checkpointing

Recompute activations (trade compute for memory)

💡 Rule of Thumb

OOM Error?

Reduce batch size by 50%

🔍 Monitoring Tools

`nvidia-smi`

`torch.cuda.memory_allocated()`

```
$ nvidia-smi
+-----------------------------------------------------------------------------+
| NVIDIA-SMI 535.104.05   Driver Version: 535.104.05   CUDA Version: 12.2     |
```

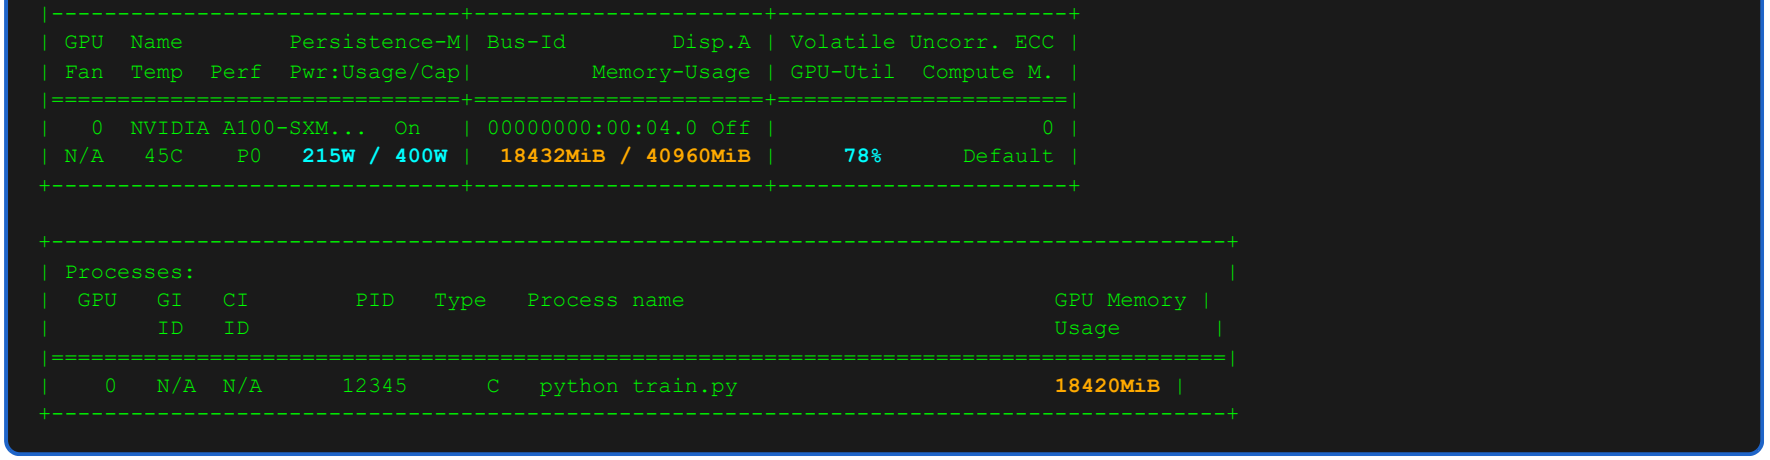
GPU Name Persistence-M	Bus-Id Disp.A	Volatile Uncorr. ECC
Fan Temp Perf Pwr:Usage/Cap	Memory-Usage	GPU-Util Compute M.
0 NVIDIA A100-SXM... On	00000000:00:04.0 Off	0
N/A 45C P0 215W / 400W	18432MiB / 40960MiB	78% Default
Processes:
GPU GI CI PID Type Process name GPU Memory
ID ID Usage
0 N/A N/A 12345 C python train.py 18420MiB