

Implementation Tips and Best Practices

Data & Training

1 Normalize Inputs

Scale features to [0,1] or standardize (mean=0, std=1)

2 Weight Initialization

Use Xavier/He initialization, avoid zeros

3 Split Data Properly

Train (70-80%), Validation (10-15%), Test (10-15%)

4 Monitor Training

Track both training and validation loss

5 Use Early Stopping

Stop when validation loss stops improving

Debugging

1 Start Simple

Begin with small network, overfit small dataset first

2 Check Gradient Flow

Verify gradients are not vanishing or exploding

3 Visualize Activations

Plot activation distributions across layers

4 Numerical Gradient Check

Verify backprop implementation with finite differences

5 Log Everything

Track loss, accuracy, learning rate, gradients



Hyperparameters



Optimization

<p>1 Learning Rate Start with 0.001-0.01, use learning rate schedules</p> <p>2 Batch Size Common: 32, 64, 128. Larger = more stable but slower</p> <p>3 Number of Epochs Start high, use early stopping to prevent overfitting</p> <p>4 Regularization Try L2 ($\lambda=0.01$), dropout ($p=0.5$), or both</p>	<p>1 Use Adam Optimizer Good default choice with adaptive learning rates</p> <p>2 Batch Normalization Normalize layer inputs for faster convergence</p> <p>3 Learning Rate Decay Reduce LR when validation loss plateaus</p> <p>4 Gradient Clipping Prevent exploding gradients (clip at norm=5)</p>
--	--

Grid/Random Search

Data Augmentation

Before Training	During Training	After Training
<ul style="list-style-type: none"> <input checked="" type="checkbox"/> Data normalized <input checked="" type="checkbox"/> Train/val/test split <input checked="" type="checkbox"/> Weights initialized <input checked="" type="checkbox"/> Loss function chosen 	<ul style="list-style-type: none"> <input checked="" type="checkbox"/> Monitor loss curves <input checked="" type="checkbox"/> Check gradient norms <input checked="" type="checkbox"/> Save best model <input checked="" type="checkbox"/> Log metrics regularly 	<ul style="list-style-type: none"> <input checked="" type="checkbox"/> Evaluate on test set <input checked="" type="checkbox"/> Analyze errors <input checked="" type="checkbox"/> Document results <input checked="" type="checkbox"/> Compare baselines

Common Pitfalls to Avoid

- ⚠️ Not shuffling data • Forgetting to normalize • Using test set for hyperparameter tuning • Setting learning rate too high • Ignoring validation metrics • Not saving checkpoints



Hands-On: TensorFlow Playground

Visualize and understand how neural networks work! Use TensorFlow Playground to observe the training process in real-time by adjusting various hyperparameters. Start with the Circle dataset and experiment with learning rates, layer structures, activation functions, and more.



Launch TensorFlow Playground