

Key Limitations of Standard RNNs



Vanishing Gradient Problem



- Gradients diminish exponentially during backpropagation
- Early time steps receive negligible weight updates
- Hinders learning from distant past information



Long-term Dependencies



Information decays over long sequences

- Difficulty capturing dependencies across long sequences
- Context from early tokens gets lost or diluted
- Limits performance on tasks requiring long-range context



Exploding Gradient Problem



Sequential Processing Bottleneck





- Gradients grow exponentially during backpropagation
- Causes numerical instability and training divergence
- Requires gradient clipping to stabilize training

Cannot parallelize ▲

- Must process sequence step-by-step sequentially
- Cannot leverage parallel computation effectively
- Slow training and inference on long sequences