

RNN

Recurrent Neural Networks
Step-by-step sequential processing



Problems

- ⚠ Vanishing gradient
- ⚠ Sequential processing limits
- ⚠ Long sequence challenges



LSTM & GRU

Addressed vanishing gradient with gates



Remaining Issues

- ⚠ Still sequential processing
- ⚠ Limited parallelization
- ⚠ Long-range dependencies



Attention Mechanism

Improved encoder-decoder models
Direct connections to all positions



Solution: Transformer Architecture

