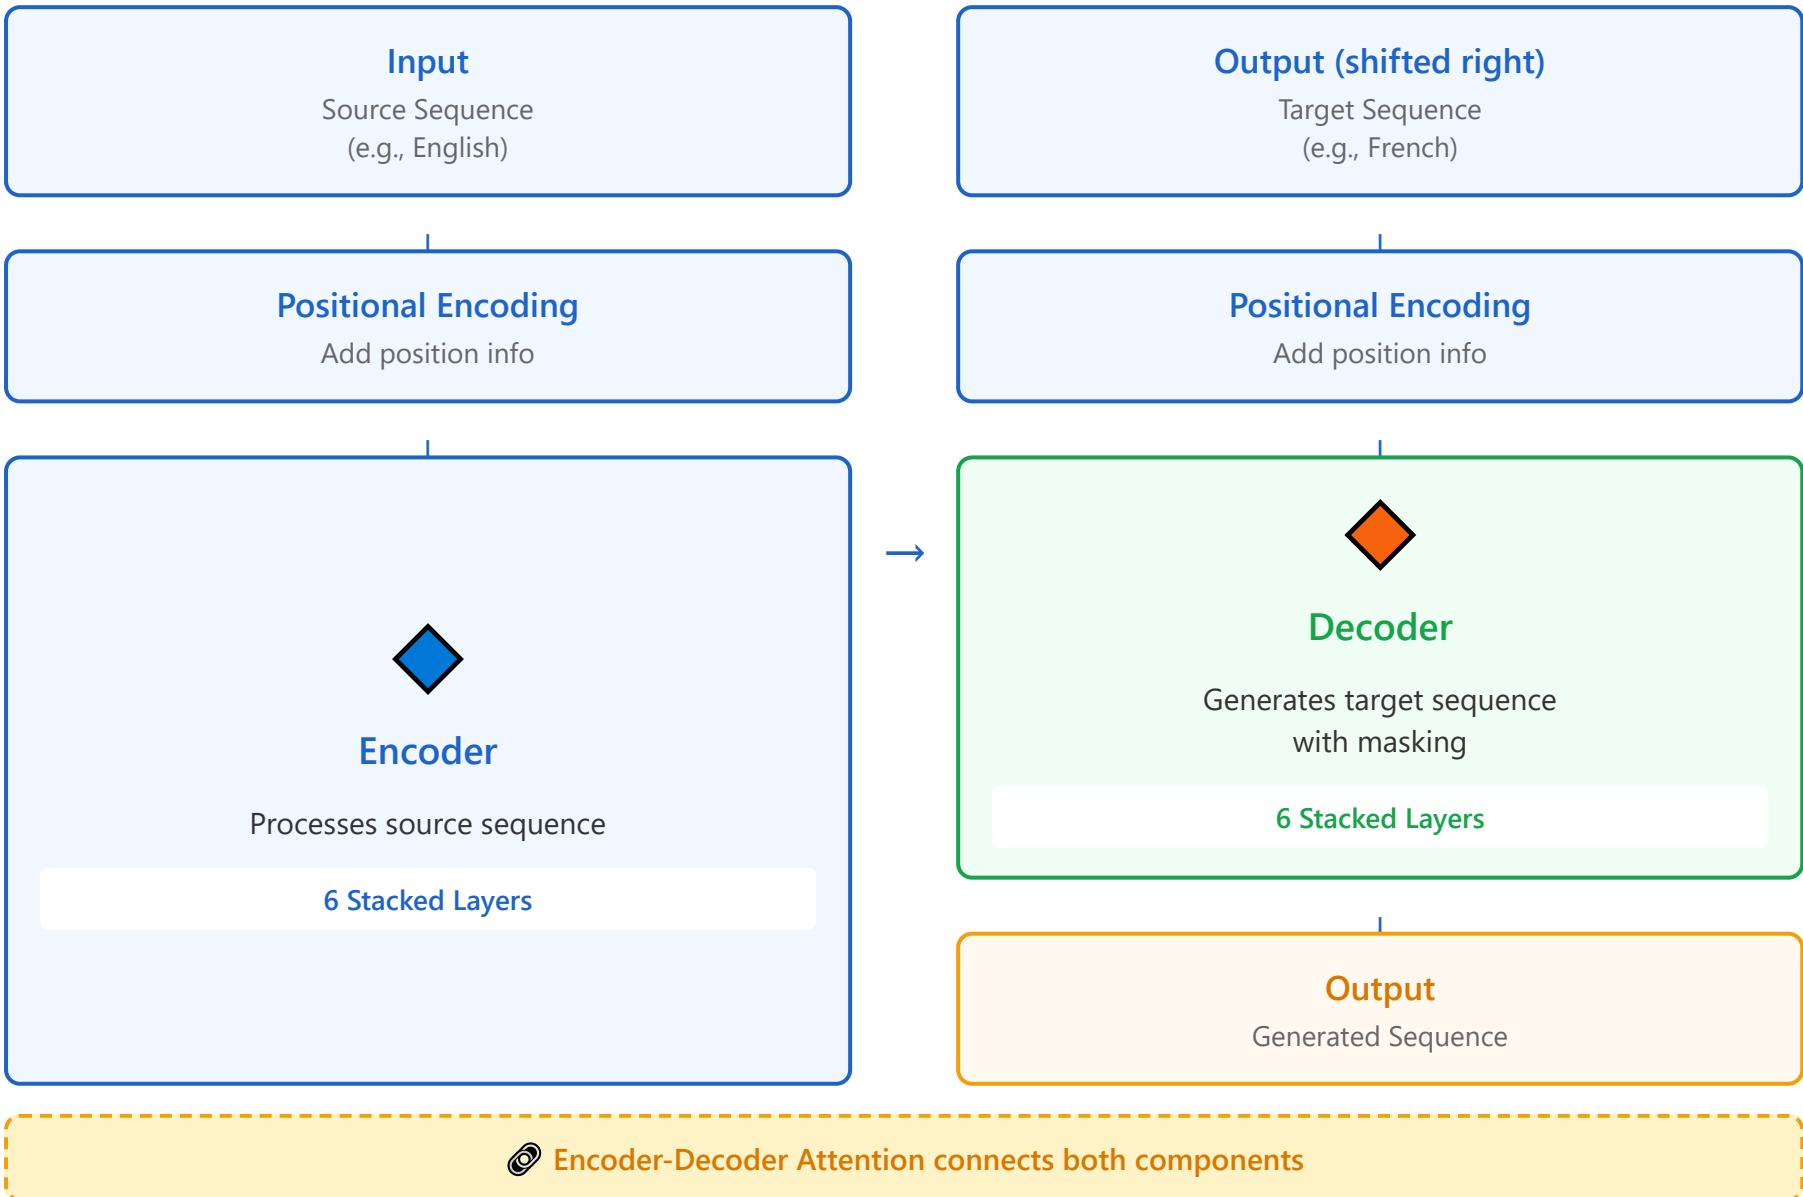


# Overall Transformer Structure

Encoder-Decoder Architecture for Sequence-to-Sequence Tasks





Identical layers  
with different parameters



## Detailed Computation Example

Encoder tasks

Translation, Summarization



Decoder uses  
**masking** for causality

Step by step process with actual dimensions ( $d_{model} = 4$ , sequence length = 8)

### Step 1: Input Embedding

**Source (English):** "I love deep learning very much!"

**Tokenization:** [I, love, deep, learning, very, much, !, <PAD>] → 8 tokens

**Embedding Matrix:** Vocabulary ×  $d_{model}$  → Each token → 4D vector

```
X_embedded (8 × 4) =
[[0.2, 0.5, -0.3, 0.8], # I
 [0.1, -0.4, 0.6, 0.3], # love
 [-0.5, 0.2, 0.4, -0.1], # deep
 [0.3, 0.7, -0.2, 0.5], # learning
 [0.4, -0.3, 0.1, 0.6], # very
 [-0.2, 0.5, 0.3, -0.4], # much
 [0.6, 0.1, -0.5, 0.2], # !
 [0.0, 0.0, 0.0, 0.0]] # <PAD>
```

### Step 2: Positional Encoding

**Add position information using sin/cos functions:**

$$PE(pos, 2i) = \sin(pos / 10000^{(2i/d_{model})})$$

$$PE(pos, 2i+1) = \cos(pos / 10000^{(2i/d_{model})})$$

```
PE (8 × 4) =
[[0.00, 1.00, 0.00, 1.00], # pos 0
 [0.84, 0.54, 0.01, 1.00], # pos 1
```

```
[0.91, -0.42, 0.02, 1.00], # pos 2  
[0.14, -0.99, 0.03, 1.00], # pos 3  
[-0.76, -0.65, 0.04, 1.00], # pos 4  
[-0.96, 0.28, 0.05, 1.00], # pos 5  
[-0.28, 0.96, 0.06, 1.00], # pos 6  
[0.66, 0.75, 0.07, 0.99]] # pos 7
```

```
X_input (8 × 4) = X_embedded + PE  
[[0.20, 1.50, -0.30, 1.80],  
 [0.94, 0.14, 0.61, 1.30],  
 [0.41, -0.22, 0.42, 0.90],  
 [0.44, -0.29, -0.17, 1.50],  
 [-0.36, -0.95, 0.14, 1.60],  
 [-1.16, 0.78, 0.35, 0.60],  
 [0.32, 1.06, -0.44, 1.20],  
 [0.66, 0.75, 0.07, 0.99]]
```

### ◆ Step 3: Encoder Self-Attention (Layer 1)

**Multi-Head Attention computation:**

**1) Linear projections ( $d_k = d_{model} / num\_heads = 4 / 2 = 2$  per head):**

```
W_Q (4 × 4), W_K (4 × 4), W_V (4 × 4)  
Q = X_input @ W_Q → (8 × 4)  
K = X_input @ W_K → (8 × 4)  
V = X_input @ W_V → (8 × 4)
```

**2) Attention scores:**

```
Scores = Q @ K^T / √d_k → (8 × 8)  
// Example scores (showing first 4×4 block):  
[[2.3, 0.8, 1.2, 1.5, ...],  
 [0.8, 2.1, 0.9, 1.3, ...],
```

```
[1.2, 0.9, 2.5, 1.1, ...],  
[1.5, 1.3, 1.1, 2.2, ...]]
```

```
Attention = softmax(Scores) → (8 × 8)  
// Softmax normalizes each row to sum to 1  
[[0.35, 0.08, 0.12, 0.15, ...], # I attends to all  
 [0.10, 0.31, 0.11, 0.16, ...], # love attends to all  
 [0.13, 0.09, 0.41, 0.12, ...], # deep attends to all  
 ...]
```

### 3) Apply attention to values:

```
Output = Attention @ V → (8 × 4)  
// Each token representation is now context-aware
```

## Step 4: Add & Norm + Feed-Forward

### 1) Residual connection and Layer Normalization:

```
X1 = LayerNorm(X_input + Attention_output) → (8 × 4)
```

### 2) Feed-Forward Network (expand to $d_{ff} = 16$ , then back to 4):

```
FFN(x) = W_2 @ ReLU(W_1 @ x + b_1) + b_2  
W_1: (4 × 16), W_2: (16 × 4)
```

```
FFN_output = FFN(X1) → (8 × 4)  
X_encoder1 = LayerNorm(X1 + FFN_output) → (8 × 4)
```

- Repeat this process for 6 encoder layers
- Final encoder output: (8 × 4)

## ◆ Step 5: Decoder Input & Masked Self-Attention

**Target (French):** "<BOS> J' aime l' apprentissage profond <PAD>"

**Shifted right:** [<BOS>, J', aime, l', apprentissage, profond, <PAD>, <PAD>]

**After embedding + positional encoding:** Y\_input ( $8 \times 4$ )

### Masked Self-Attention:

Mask prevents attending to future tokens:

```
[[1, 0, 0, 0, 0, 0, 0, 0], # <BOS> only sees itself  
[1, 1, 0, 0, 0, 0, 0, 0], # J' sees <BOS>, J'  
[1, 1, 1, 0, 0, 0, 0, 0], # aime sees up to aime  
[1, 1, 1, 1, 0, 0, 0, 0], # l' sees up to l'  
[1, 1, 1, 1, 1, 0, 0, 0], # apprentissage  
[1, 1, 1, 1, 1, 1, 0, 0], # profond  
[1, 1, 1, 1, 1, 1, 1, 0], # <PAD>  
[1, 1, 1, 1, 1, 1, 1, 1]] # <PAD>
```

Masked\_Attention\_output → ( $8 \times 4$ )

## ⌚ Step 6: Encoder-Decoder Cross-Attention

### Query from decoder, Key & Value from encoder:

```
Q = Decoder_output @ W_Q → (8 × 4) # from decoder  
K = Encoder_output @ W_K → (8 × 4) # from encoder  
V = Encoder_output @ W_V → (8 × 4) # from encoder
```

```
Cross_Attention = softmax(Q @ K^T / √ d_k) @ V  
// Decoder attends to ALL encoder positions  
// This allows decoder to "look at" the source sentence
```

Cross\_Attention\_output → (8 × 4)

## ⬆ Step 7: Final Output & Prediction

**After 6 decoder layers:**

Decoder\_final\_output → (8 × 4)

**Linear projection to vocabulary size (e.g., 10,000):**

Logits = Decoder\_output @ W\_output → (8 × 10000)

Probabilities = softmax(Logits) → (8 × 10000)

**For each position, pick highest probability token:**

Position 0:  $P("J'") = 0.92 \rightarrow \text{output } "J'"$

Position 1:  $P("aime") = 0.87 \rightarrow \text{output } "aime"$

Position 2:  $P("l'") = 0.91 \rightarrow \text{output } "l'"$

Position 3:  $P("apprentissage") = 0.84 \rightarrow \text{output } "apprentissage"$

Position 4:  $P("profond") = 0.89 \rightarrow \text{output } "profond"$

Position 5:  $P("<\text{EOS}>") = 0.95 \rightarrow \text{output } "<\text{EOS}>" \text{ (stop)}$

**Final Translation: "J' aime l' apprentissage profond"**



## Key Dimensional Flow Summary

Input tokens (8) → Embedding (8×4) → +PE (8×4) →

Encoder layers (8×4 maintained) → Encoder output (8×4) →

Decoder input ( $8 \times 4$ ) → Decoder layers ( $8 \times 4$  maintained) →

Linear projection ( $8 \times 4 \rightarrow 8 \times 10000$ ) → Softmax → Predictions ( $8 \times 10000$ )