

Summary and Next Steps

Key Concepts

 **Self-Attention Mechanism**

 **Multi-Head Attention**

 **Positional Encoding**

Advantages

 **Parallelization** - Fast training

 **Long-range dependencies** - Global context

 **Flexibility** - Multiple domains

Challenges

 Quadratic complexity $O(n^2)$

 Large memory requirements

Next Steps

1 Next Lecture

Advanced topics: Sparse attention, Linear attention

2 Efficient Transformers

Explore Reformer, Linformer, Performer

3 Practice

Implement a mini-Transformer from scratch

4 Explore Libraries

Hugging Face Transformers library

Recommended Reading

"Attention Is All You Need" paper

Vaswani et al., 2017