

Hands-on: Image Representation Learning with SimCLR

PyTorch Implementation Workflow

1 Dataset Setup

Load unlabeled dataset for pretraining



CIFAR-10 / STL-10

2 Implement SimCLR

Build encoder & projection head

PyTorch



3 Data Augmentation

Setup augmentation pipeline

Crop + Color + Blur



4 Train Encoder

Optimize with contrastive loss

NT-Xent Loss



5 Linear Probe

Evaluate on labeled subset

Frozen Features



6 Visualize Embeddings

Plot learned representations

t-SNE



Final Step: Compare self-supervised results with supervised baseline performance



Example Data and Calculation Process



1. Input Images & Augmentation

Image	Size	Augmentation
Cat (Original)	32×32×3	-
Cat View 1	32×32×3	Random Crop + Flip
Cat View 2	32×32×3	Color Jitter + Blur

Key Point: Generate two different views from the same image to form a positive pair



2. Encoder & Projection Head

Architecture:

Input (3×32×32) → ResNet-18 → h (512-dim) → MLP → z (128-dim)

Example embedding vectors

$z_1 = [0.23, -0.45, 0.67, \dots, 0.12]$ # 128-dim

$z_2 = [0.21, -0.43, 0.65, \dots, 0.15]$ # 128-dim (positive)

$z_3 = [-0.89, 0.34, -0.12, \dots, 0.56]$ # 128-dim (negative)

Key Point: Calculate cosine similarity after L2 normalization



3. NT-Xent Loss Calculation (Batch Size N=4, Total 8 Samples)

NT-Xent Loss Formula:

$$\ell(i, j) = -\log[\exp(\text{sim}(z_i, z_j)/\tau) / \sum \exp(\text{sim}(z_i, z_k)/\tau)]$$

τ (temperature) = 0.5

Step 1: Calculate Similarity Matrix (8×8)

```
sim(z1, z2) = 0.92 // positive pair  
sim(z1, z3) = 0.15 // negative  
sim(z1, z4) = -0.23 // negative  
sim(z1, z5) = 0.08 // negative  
...  
...
```

Step 2: Temperature Scaling ($\tau=0.5$)

```
sim(z1, z2) / τ = 0.92/0.5 = 1.84  
sim(z1, z3) / τ = 0.15/0.5 = 0.30  
sim(z1, z4) / τ = -0.23/0.5 = -0.46  
...
```

Step 3: Exponential Calculation

```
exp(1.84) = 6.297 // positive (numerator)  
exp(0.30) = 1.350 // negative  
exp(-0.46) = 0.631 // negative  
exp(0.16) = 1.174 // negative  
...  
Denominator Sum = 6.297 + 1.350 + 0.631 + 1.174 + ... = 12.845
```

Step 4: Loss Calculation

```
l(z1, z2) = -log(6.297 / 12.845)  
l(z1, z2) = -log(0.490)  
l(z1, z2) = 0.713
```

Final Batch Loss (Average)

Loss = 0.713

 **Interpretation:** Lower loss indicates that the model has learned good representations where positive pairs are close and negative pairs are far apart.

Training Results Example

⚡ Self-Supervised Pretraining

Epoch	Loss	Time
1	2.456	8 min
50	1.234	-
100	0.876	-
200	0.623	26 hours

Settings: CIFAR-10 unlabeled, batch=256, ResNet-18

📊 Linear Probe Evaluation (CIFAR-10 Test)

Method	Accuracy	Training Data
Random Init	23.4%	5,000 labels
SimCLR (ours)	78.6%	5,000 labels
Supervised	91.2%	50,000 labels

Self-Supervised Improvement

+55.2%

💻 PyTorch Implementation (NT-Xent Loss)

```
# SimCLR NT-Xent Loss Implementation
def nt_xent_loss(z_i, z_j, temperature=0.5):
    """
    Args:
        z_i, z_j: [batch_size, embedding_dim] - positive pair
    Returns:
        loss: scalar tensor
    """
    batch_size = z_i.shape[0]

    # L2 normalization
    z_i = F.normalize(z_i, dim=1)
    z_j = F.normalize(z_j, dim=1)
```

```
# Concatenate: [2*batch_size, embedding_dim]
z = torch.cat([z_i, z_j], dim=0)

# Similarity matrix: [2N, 2N]
similarity_matrix = torch.mm(z, z.T) / temperature

# Mask to exclude self-similarity
mask = torch.eye(2 * batch_size, dtype=torch.bool)
similarity_matrix = similarity_matrix.masked_fill(mask, -9e15)

# Positive pairs: (i, N+i) and (N+i, i)
positives = torch.cat([
    torch.diag(similarity_matrix, batch_size),
    torch.diag(similarity_matrix, -batch_size)
])

# NT-Xent loss calculation
nominator = torch.exp(positives)
denominator = torch.sum(torch.exp(similarity_matrix), dim=1)
loss = -torch.log(nominator / denominator)

return loss.mean()
```



Key Insight: Self-supervised learning can learn powerful feature representations without labels, achieving high performance with only a small amount of labeled data.