

Attention Mechanisms as Explanations: Transformer Visualization

Self-Attention Matrix (BERT Example)

	This	movie	is	great	!	[SEP]
This	0.42	0.18	0.08	0.15	0.05	0.02
movie	0.20	0.35	0.10	0.22	0.08	0.05
is	0.08	0.15	0.18	0.45	0.12	0.02
great	0.10	0.25	0.18	0.38	0.07	0.02
!	0.05	0.08	0.05	0.28	0.48	0.06
[SEP]	0.12	0.10	0.08	0.15	0.10	0.25

■ High (>0.35) ■ Med-High ■ Medium ■ Low ■ Very Low

Token Attention Example: "movie" attending to:



Self-Attention

Shows token-to-token importance in context

- Query-Key-Value mechanism
- Context-aware weights
- Bidirectional relationships



Multi-Head Attention

Different heads capture different aspects

- Parallel attention layers
- Diverse relationship types
- Aggregate for full picture



Attention Methods

- **Rollout:** Aggregate across layers
- **Flow:** Trace information path
- **Gradient:** Combine with gradients



Important Limitation

Attention ≠ Explanation

Recent research shows attention weights don't always reflect true

This movie is great !

model reasoning. Combine with gradient-based methods for better insights.