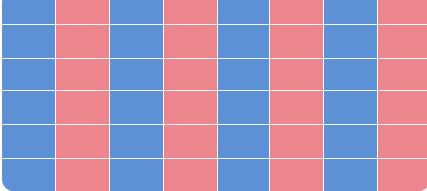


Gradient-based Methods: Computing Feature Sensitivity

Vanilla Gradients



$\partial y / \partial x$

Basic sensitivity, noisy

Integrated Gradients


$$\int_0^1 \nabla f(x^- + \alpha(x-x)) d\alpha$$

Path integral, smooth

SmoothGrad


$$E[\nabla f(x + N(0, \sigma^2))]$$

Averaged over noise

Grad-CAM


$$\text{ReLU}(\Sigma \alpha_k A_k)$$

Class activation map

Method Comparison

Vanilla Fast, noisy ✗

Integrated Smooth, slow ✓

SmoothGrad Less noise ✓

Grad-CAM CNN specific ✓

Strengths

- ✓ Model-agnostic
- ✓ Fast computation
- ✓ Differentiable

Limitations

- ✗ Saturation issue
- ✗ Noise in vanilla
- ✗ Requires gradients

PyTorch Implementation

Vanilla Gradients

```
x.requires_grad_()
output = model(x)
output.backward()
gradients = x.grad
```

Integrated Gradients

```
# Path from baseline to input
alphas = torch.linspace(0, 1, steps=50)
for alpha in alphas:
    x_step = baseline + alpha * (x - baseline)
    # Compute gradients at each step
```

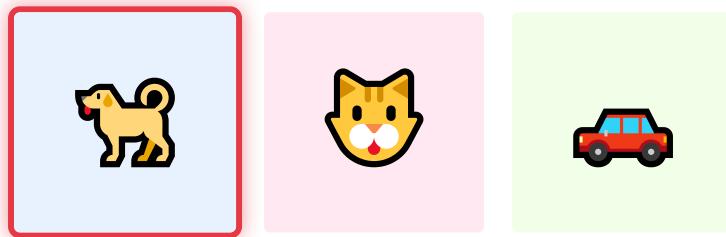
SmoothGrad

```
# Average over noisy samples
for _ in range(n_samples):
```

```
noise = torch.randn_like(x) * sigma  
grads += compute_grad(x + noise)  
smooth_grad = grads / n_samples
```

🔍 Interactive CAM Visualization Demo

📷 Select Image



🔧 Method Selection

Grad-CAM

Integrated Gradients

SmoothGrad

Vanilla Gradients

🎨 Heatmap Settings

Opacity

60%

Focus Intensity

70%

CAM Visualization on Image





High → Low Activation